
SelfBound: A Diagnostic Benchmark for LLM Agent Self-Knowledge under Noisy Feedback

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Users learn where each LLM agent is reliable; the agent itself does not. We ask
2 a single mechanistic question: at which sub-capability does this self-modeling
3 break down? Three candidates—H1 output channel, H2 prior metacognition,
4 H3 noise-robust online inference—are conflated by standard benchmarks. We
5 build **SelfBound**, a diagnostic benchmark whose five-level oracle ladder (L0 raw
6 confidence $\rightarrow$ L4 full-statistic oracle) lets each adjacent comparison rule out one
7 candidate, scored with **Capability Boundary Fidelity (CBF)**, a per-domain metric
8 distinct from per-query calibration (ECE). The ladder yields a clean deduction
9 across 15 models: **H1 is ruled out** (oracle CBF ≥ 0.99 on every model); **H2**
10 **is partial and family-split** rather than the bottleneck (GPT models reach 0.795–
11 0.859 zero-shot; Claude models fall below NoMemory); **H3 is the locus**—first-
12 order single-rater learners cluster near NoMemory, while explicit *second-order*
13 estimators (asymmetric Dawid–Skene, GLAD) substantially exceed the L1 prior,
14 with GLAD beating L1 on all six frontier models ($p = 0.031$). The deduction
15 repeats across task families (12/15 sign-flip rate on extension tasks), implying
16 per-task boundary modeling. We name this failure mode **Self-Boundary Collapse**
17 and isolate noise-robust second-order online inference as the open problem; the
18 ladder, CBF, cross-family user simulator, and trace dataset are released as the
19 diagnostic infrastructure.

20 1 Introduction

21 LLM agents are deployed in long-horizon, multi-turn settings (customer service, coding, research,
22 tool-use): users develop an empirical sense of where each agent is reliable, but the agent itself retains
23 no analogous self-model—each session starts without prior calibration and discards its own. The
24 asymmetry is documented, but the underlying *capability gap* is poorly understood: we still do not
25 know *where*, mechanistically, agent self-modeling actually fails.

26 **Three competing hypotheses about where self-knowledge fails.** Existing evaluations conflate
27 three structurally distinct bottlenecks:

- 28 • **H1 (Output channel).** The model cannot numerically format and emit per-domain statistics.
- 29 • **H2 (Prior metacognition).** The model lacks an accurate prior over its own per-domain abilities,
30 so any self-assessment must be *learned* from interaction.
- 31 • **H3 (Noise-robust online inference).** The model can report correct priors but cannot reliably
32 update them from *noisy* implicit user feedback.

33 These imply very different research directions—H1 motivates fine-tuning on numeric outputs; H2
34 motivates pre-training-time metacognitive curricula; H3 motivates noise-robust online estimation—
35 but Barkan et al. [3], Hu et al. [16], calibration research [13, 18, 45], and existing memory systems [26,

33, 35, 41, 47] all probe the symptom without isolating the locus. We name the resulting failure mode—LLM agents repeatedly receiving feedback yet failing to translate it into a stable, domain-level self-model—**Self-Boundary Collapse**, and ask: at which level of the H1/H2/H3 ladder does it actually occur?

Benchmark design and the CBF metric. We answer the question with SelfBound, whose central design choice is an oracle-bracketed *five-level baseline ladder*—L0 raw confidence (NoMemory), L1 zero-shot per-domain prior, L2 noisy-feedback learners (LLM-self-assessment, classical recalibration, EM-joint, second-order Dawid–Skene/GLAD), L3 binary-label oracle, L4 full-statistic oracle. Each adjacent comparison rules out one hypothesis: L4–L3 isolates the output channel (H1), L3–L1 isolates the prior (H2), L3–L2 isolates online inference (H3). The protocol pairs each agent with a domain-heterogeneous simulated user (75% reliable in-expertise, 46% out, with cross-family routing to mitigate a dual-role bias we report quantitatively in §4.3) over 50-turn sessions, and scores the resulting per-domain self-model with **Capability Boundary Fidelity (CBF)**, a metric we introduce that complements per-query calibration measures such as ECE.

Localization across the L0–L4 ladder. Across 15 models from 7 families and 47K+ interactions:

- **H1 is ruled out.** GT-SelfStats reaches $\text{CBF} \geq 0.99$ on every model: when given the ground-truth per-domain statistics in-prompt, the output channel reproduces them.
- **H2 is partial and family-split, not the bottleneck.** The zero-shot prior is the strongest non-oracle baseline on GPT (0.795 on GPT-5.4, 0.859 on GPT-5.5) but falls below NoMemory on all three Claude models (0.460–0.482 vs. 0.557–0.603); frontier mean 0.640. Prior knowledge exists for some families and is absent for others, but never reaches the oracle ceiling.
- **H3 is the locus.** First-order single-rater learners (UF-PDC, Platt, HistBin, Bayes, simple EM-Joint) cluster near NoMemory on the frontier (0.628–0.634), unable to recover the per-domain truth from noisy feedback. Explicit *second-order* estimators, asymmetric Dawid–Skene (frontier mean 0.755) and GLAD (0.712), substantially exceed the L1 prior, with GLAD beating L1 on all six frontier models (Wilcoxon two-sided $p = 0.031$). The L2→L3 gap of ~ 0.30 CBF remains, but it is structurally addressable, not a fundamental wall.

2 Related Work

LLM self-knowledge and calibration benchmarks. A growing line of work probes whether LLMs can represent or report their own correctness on a single query. Internal-representation studies show that latent confidence signals exist [4, 17, 18], while recent benchmarks evaluate metacognition behaviorally and find that frontier models exhibit only partial, context-dependent self-monitoring [1, 30]. Table 1 positions the most closely related benchmarks against SelfBound along four design axes: feedback channel, horizon, whether self-knowledge is accumulated, and whether the protocol forces second-order user-reliability inference.

Table 1: Related-benchmark positioning along four design axes. SelfBound is the only benchmark with a noisy interactive feedback channel that the agent must accumulate into a per-domain self-model under second-order user-reliability uncertainty.

Benchmark	Feedback	Horizon	Self-knowledge	2nd-order
SAD [22], SelfAware [51]	none	single-shot	static probe	—
Barkan [3]	clean step signals	multi-step exec	per-task	—
MemoryAgentBench [16]	implicit (mem ops)	long-horizon	exogenous mem only	—
AgentBoard [31], τ -bench [49]	task reward only	multi-turn	none	—
Knowledge-boundary [24, 44]	none	single-shot	static	—
SelfBound (ours)	noisy NL, 75/46 asym	50-turn session	per-domain accumulated	✓

(i) *Single-shot self-knowledge probing.* SAD [22] and SelfAware [51] score self-knowledge per query without temporal dynamics; their findings supply the “L1 prior” anchor of our ladder. SelfBound uses the L1→L3 contrast to ask whether *noisy interaction* improves on that prior—finding it does not.

(ii) *Multi-step agent overconfidence.* Barkan et al. [3] measure whether 13 frontier LLMs can predict their own success on BigCodeBench and SWE-Bench under multi-step tool execution and find all models overconfident with overconfidence *worsening* during execution. Their setting gives clean




Figure 1: **The L0–L4 baseline ladder localizes where agent self-knowledge fails.** L4≈L3≈1.0 rules out the output channel (H1); the L2→L3 gap of ~0.30 isolates noise-robust online inference (H3) as the bottleneck. **(A)** Session protocol: 50 sessions × 50 turns over 20 MMLU domains (green=correct, red=wrong). **(B)** Single turn: the agent answers with confidence; a cross-family simulated user gives domain-dependent noisy feedback ($r_{in} \approx 75\%$ in-expertise, $r_{out} \approx 46\%$ out); a decoder extracts the labeled signal that updates the per-domain self-model. **(C)** The five-level diagnostic ladder: L0 raw confidence (NoMemory) → L1 zero-shot prior (ZS-SK) → L2 noisy-feedback learners (LSA, EM-Joint, DS-asym, GLAD) → L3 binary-label oracle (OraclePDC) → L4 full-statistic oracle (GT-SelfStats). Adjacent comparisons isolate each candidate failure: L4–L3 isolates H1, L3–L1 isolates H2, L3–L2 isolates H3.

per-step success/failure signals; ours gives noisy domain-dependent natural-language feedback, where first-order methods cluster at the L2 floor. Barkan establishes that overconfidence persists under clean traces; we localize *where* the failure lives when feedback itself is noisy, via the L0–L4 oracle-bracketing ladder.

(iii) *Multi-turn agent benchmarks without self-knowledge.* AgentBoard [31] and τ -bench [49] evaluate task completion under multi-turn interaction but do not measure self-knowledge; their reward signal is task success, not self-estimate accuracy. SelfBound holds the underlying task simple (MCQ) and varies only the information-access regime, so the diagnostic ladder isolates self-knowledge mechanisms these benchmarks cannot disentangle from task-execution confounds.

Static knowledge-boundary benchmarks [12, 15, 24, 27, 28, 44, 50] share the same structural limitation: they evaluate *a priori* boundary awareness, not whether the agent’s self-model can be accumulated under noisy interactive turns (Appendix R).

Agent memory and learning from implicit feedback. Existing memory systems—MemGPT [35], Mem0 [33], RMM [41], A-MEM [47], MemOS [26], MEM1 [54]—and the recent MemoryAgent-Bench [16] all store *exogenous* state (user facts, dialogue, trajectories). None model the agent’s *own* per-domain reliability; we introduce this as a fifth memory competency. On the feedback-learning side, RLHF [34] and DPO [37] update weights at training time assuming reliable annotators; Liu et al. [29] establish that implicit dialogue feedback is noisy and only partially useful as a distillation signal; Leng et al. [23] show RLHF itself induces systematic overconfidence. SelfBound studies the orthogonal regime: *inference-time* accumulation of *noisy, domain-dependent* feedback into a per-domain self-model, with no parameter updates. The unique twist is asymmetric reliability (75% in-expertise vs. 46% out), which turns boundary learning into a *second-order* problem.

99 **Selective prediction and task-specific calibration.** Per-query selective-prediction methods [2, 9,
100 19, 20, 25, 42, 46] treat each prediction in isolation. SelfBound instead enables *historical* selective
101 prediction: abstention based on persistent per-domain self-model statistics. Independently, prior
102 work has noted that calibration does not transfer across tasks [6, 32, 39]; our experiments quan-
103 tify this systematically across MMLU, GSM8K, HumanEval, and BFCL. The closest concurrent
104 benchmarks [5, 10, 11, 14, 53] either assume perfect feedback, require weight updates, or focus on
105 single-agent dynamics; none combine multi-task evaluation, noisy domain-dependent feedback, and
106 a quantitative boundary-fidelity metric over extended horizons.

107 3 SelfBound: Probing Self-Knowledge Memory under Noisy Feedback

108 SelfBound evaluates a single capability: given a stream of (q_t, a_t, c_t, f_t) tuples over T turns—a
109 question q_t from a hidden domain d_t , the agent’s answer a_t with confidence $c_t \in [0, 1]$, and a
110 natural-language user feedback f_t —can the agent build a function $\hat{p} : \mathcal{D} \rightarrow [0, 1]$ that closely
111 approximates its true per-domain accuracy $p^*(d)$? The benchmark instantiates this question via
112 a 50×50 protocol (50 sessions $\times$ 50 turns per session = 2,500 turns per model) drawing from 20
113 MMLU subjects [15] partitioned per model into Strong/Medium/Weak tiers relative to that model’s
114 own accuracy distribution. The user is simulated by a cross-family LLM (avoiding the dual-role bias
115 of Appendix M) parameterised by an expertise set $\mathcal{E} \subset \mathcal{D}$ with empirically calibrated reliabilities of
116 $\sim 75\%$ in-expertise and $\sim 46\%$ out, and we additionally construct three extension tasks—GSM8K [8],
117 HumanEval [7], and BFCL tool-calling [48]—collectively called the *extension suite*—to test cross-
118 task generalisation (Appendix P). Full task formalisation, session protocol, question-pool partitioning,
119 and persona parameterisation are deferred to Appendix B.

120 **Capability Boundary Fidelity (CBF).** Our central metric scores how closely the agent’s self-
121 assessed per-domain accuracy tracks the truth:

$$\text{CBF} = 1 - \frac{1}{|\mathcal{D}|} \sum_{d \in \mathcal{D}} |\hat{p}(d) - p^*(d)|, \quad \text{CBF} \in [0, 1]. \quad (1)$$

122 CBF is computed once per session at the end of the 50-turn protocol. CBF is complementary to
123 per-query calibration metrics (ECE/Brier): ECE operates on (c_t, y_t) pairs while CBF operates on
124 $(\hat{p}(d), p^*(d))$ pairs, so neither implies the other—we verify this empirically in Appendix B. The
125 trivial estimator $\hat{p}(d) = 0.5$ attains $\text{CBF} = 1 - \mathbb{E}_d[|0.5 - p^*(d)|]$, which is high for weak models
126 clustered near 50% and low for strong models with widely varying accuracy across domains; this
127 asymmetry is exactly what makes the metric informative for our diagnostic.

128 **Baseline-adjusted CBF: the primary cross-model metric.** To remove the trivial-estimator artifact
129 when comparing methods *across* models, we define a normalized variant that measures the fraction
130 of the oracle gap a method closes relative to the constant baseline:

$$\widetilde{\text{CBF}}(m, \text{method}) = \frac{\text{CBF}_{\text{method}}(m) - \text{CBF}_{\text{NoMem}}(m)}{\text{CBF}_{\text{Oracle}}(m) - \text{CBF}_{\text{NoMem}}(m)}, \quad (2)$$

131 where $\text{CBF}_{\text{NoMem}}(m)$ uses the model’s verbalized confidence (which empirically clusters near 0.5 on
132 the trivial domains) as the constant-baseline anchor, and $\text{CBF}_{\text{Oracle}}(m) \geq 0.99$ is the L3 OraclePDC
133 ceiling. Values are interpretable as “fraction of L0→L3 gap closed”: $\widetilde{\text{CBF}} = 1$ matches oracle, 0
134 ties the trivial baseline, < 0 underperforms it. We report CBF as the primary metric for cross-model
135 means and ladder-headline comparisons (§4.2), and use raw CBF only as a within-model diagnostic.

136 **Auxiliary metrics and domain structure.** We additionally report Hallucination Rate, ECE, and
137 Brier as auxiliary calibration metrics; their definitions follow Guo et al. [13] and are given in
138 Appendix B. CBF requires a multi-domain task: for single-domain extension tasks (GSM8K, Hu-
139 manEval, BFCL) we report only ECE and confidence–accuracy gap, and the cross-task analysis in
140 Appendix P uses these standard metrics.

141 **Cross-model comparability.** The 20 MMLU subjects are partitioned into Strong/Medium/Weak
142 tiers *relative to each model’s own held-out validation accuracy* so that every model faces a non-trivial
143 cross-domain spread. CBF is therefore a within-model diagnostic; cross-model means reported

throughout (e.g., 0.685 for ZS-SK across all 15 models in Table 2) summarize per-model values that were each computed against the model’s own $p^*(d)$, and are appropriate for comparing *methods within the diagnostic ladder*, not for ranking models on absolute capability. We expand on this design choice and confirm headline conclusions are unchanged under an alternative absolute-cutoff partition in Appendix B.

Baselines: a five-level diagnostic ladder. We evaluate ten baselines organized as a ladder of *information access*, designed so that comparing adjacent levels isolates a specific sub-capability:

- **L0 (raw): NoMemory**—verbalized confidence with no memory, the stateless anchor.
- **L1 (prior only): ZS-SK** (ZeroShot Self-Knowledge)—the agent is asked, with no session history and no feedback, to estimate its own per-domain accuracy directly; this measures *prior metacognition*.
- **L2 (learning from noisy feedback): LSA** (LLMSelfAssess; in-context reasoning over session history), *UF-PDC* (LLM-decoded per-domain counting), *Platt-Online* [36], *HistBin-Online* [52], *Bayes-Domain* (Beta posterior per domain), and *EM-Joint* (joint estimation of $\hat{p}(d)$ and $\hat{r}(d)$). These differ in their inference machinery but all consume the same noisy feedback stream.
- **L3 (binary-label oracle): OraclePDC**—per-turn ground-truth correctness labels, with the agent computing per-domain running averages. Tests whether the agent could build a self-model *if feedback were clean*.
- **L4 (format-channel sanity ceiling): GT-SelfStats**—the agent is given $p^*(d)$ directly in the prompt and asked to emit them in the structured response format. This is a sanity check that the L2 gap cannot be attributed to an inability to format and emit per-domain statistics. We do not claim it tests verbalization-of-self-knowledge.

All feedback decoding (L2 + UF-PDC) uses an LLM-based decoder rather than rule-based pattern matching, more aligned with the “self-knowledge boundary” setting and avoiding rule-classifier coverage as a confound. Full per-baseline algorithms are in Appendix C.

4 Diagnostic Results: Where Does Agent Self-Knowledge Fail?

4.1 Setup

We evaluate fifteen models from seven families spanning 42–90% MMLU accuracy (Appendix S). Each model–baseline pair runs 2,500 turns (50 sessions $\times$ 50 turns under the 50 $\times$ 50 protocol); user personas are simulated by a cross-family LLM with three expertise domains¹; answer extraction uses a lenient multi-strategy parser tolerant of free-form responses to avoid truncating reasoning models. Feedback decoding for all L2 baselines uses an LLM-based decoder. Open-weights models (Mistral-7B, Llama-3-8B) are served via vLLM from pinned snapshots and serve as the reproducibility anchor; API access dates, simulator-routing rules, and the parser specification are deferred to Appendix S.

4.2 Main Results

Three diagnostic patterns dominate Table 2: (i) algorithmic L2 methods collapse to NoMem on every frontier model; (ii) the zero-shot prior splits cleanly by family (GPT wins, Claude is inflated); (iii) only second-order estimators (DS-asym, GLAD) and in-context LSA exceed the L1 prior on the frontier. We unpack each below.

Algorithmic-L2 collapse on new frontier rows, broken cleanly by ZS/LSA. For both newly added frontier models, the five *algorithmic* L2 baselines (UF-PDC, Platt, HistBin, Bayes, EM-Joint) cluster near NoMem (GPT-5.5: 0.601–0.612, all $\sim$ 0.03–0.04 above NoMem 0.570; Claude-Opus-4-7: 0.570–0.593, all $\sim$ 0.01–0.04 above NoMem 0.557). With the 50 $\times$ 50 ZS/LSA values available, both LSA and the L1 prior break this collapse decisively on GPT-5.5 (LSA 0.902, ZS-SK 0.859: LSA is +0.290 over best algorithmic L2 0.612, ZS is +0.247 over the same) and on Claude-Opus-4-7 (LSA 0.641 vs. NoMem 0.557, +0.084). The pattern under the 50 $\times$ 50 protocol: algorithmic L2

¹ Frontier-set scope: 6 model endpoints (3 GPT, 3 Claude): GPT-5.4, GPT-5.4-mini, GPT-5.5, Claude-Opus-4-6, Claude-Sonnet-4-6, Claude-Opus-4-7. Main-table within-tier means, MMLU-Pro and BBH contamination controls (Appendix E, E), and Dawid–Skene/GLAD cross-model analysis all use this $n=6$ scope. $n=6$ BBH and MMLU-Pro CBF means are 0.811 and 0.849 (vs. MMLU ZS 0.640).

Table 2: Main results: end-of-session CBF across 15 models on the five-level baseline ladder. Higher is better. **Bold** marks the best non-oracle (L0–L2) baseline per row. All 15 models use the **50 sessions** $\times$ **50 turns** protocol.

Model (Acc)	L0	L1	L2 (learn from noisy feedback)							Oracles	
	NoMem	ZS-SK	LSA	UF-PDC	Platt	HistBin	Bayes	EM	Best	OraPDC	GT-Stats
Mistral-7B (42%)	0.729	0.729	0.729	0.720	0.717	0.717	0.731	0.732	EM	1.000	1.000
Llama-3-8B (55%)	0.720	0.711	0.549	0.715	0.724	0.724	0.730	0.731	EM	1.000	1.000
GLM-5 (55%)	0.700	0.700	0.700	0.689	0.696	0.696	0.711	0.714	EM	1.000	1.000
DeepSeek-V3.2 (62%)	0.681	0.673	0.721	0.681	0.683	0.683	0.696	0.699	LSA	1.000	1.000
Qwen-Turbo (64%)	0.708	0.730	0.703	0.697	0.712	0.712	0.724	0.728	ZS	1.000	1.000
Qwen2.5-7B (65%)	0.744	0.663	0.701	0.738	0.745	0.745	0.755	0.757	EM	1.000	1.000
Qwen-Max (71%)	0.704	0.766	0.781	0.660	0.684	0.684	0.702	0.705	LSA	1.000	1.000
Qwen-Plus (79%)	0.683	0.735	0.435	0.686	0.695	0.695	0.700	0.701	ZS	1.000	1.000
Qwen3.5-27B (79%)	0.731	0.731	0.731	0.712	0.720	0.720	0.740	0.746	EM	1.000	1.000
<i>Frontier closed-source models:</i>											
GPT-5.4-mini (72%)	0.687	0.769	0.763	0.687	0.701	0.701	0.706	0.706	ZS	1.000	0.995
GPT-5.4 (82%)	0.649	0.795	0.781	0.668	0.675	0.675	0.674	0.677	ZS	1.000	0.994
GPT-5.5 (94%)	0.570	0.859	0.902	0.601	0.612	0.612	0.603	0.605	LSA	1.000	0.992
Claude-Opus-4-7 (69%)	0.557	0.482	0.641	0.593	0.577	0.577	0.571	0.570	LSA	1.000	0.992
Claude-Sonnet-4-6 (86%)	0.603	0.472	0.720	0.640	0.617	0.617	0.613	0.609	LSA	1.000	0.993
Claude-Opus-4-6 (90%)	0.594	0.460	0.680	0.618	0.597	0.597	0.603	0.604	LSA	1.000	0.993
<i>Within-tier means:</i>											
Frontier-v4 (n=6; all columns)	0.610	0.640	0.748	0.634	0.630	0.630	0.628	0.628	LSA	1.000	0.993
Strong-mid (n=3)	0.706	0.744	0.649	0.686	0.700	0.700	0.714	0.717	ZS	1.000	1.000
Mid (n=3)	0.711	0.689	0.708	0.705	0.713	0.713	0.725	0.728	EM	1.000	1.000
Weak (n=3)	0.716	0.713	0.659	0.708	0.712	0.712	0.724	0.726	EM	1.000	1.000

190 methods fail to surpass NoMem on every frontier agent; LSA cleanly exceeds NoMem on all 6
191 frontier rows; ZS exceeds NoMem on all 3 GPT rows but falls below NoMem on all 3 Claude rows. In
192 particular, GPT-5.5 is now decisively dominated by LSA (0.902, the benchmark’s highest non-oracle
193 CBF) followed by ZS (0.859); both substantially exceed every algorithmic L2 method (≤ 0.612) and
194 the prior 10×50 anomaly where LSA appeared catastrophically below NoMem (0.623) does not
195 reproduce. Held-out evaluation and second-order estimators (Dawid–Skene, GLAD) are reported on
196 the 4 original frontier models; the MMLU-Pro and BBH contamination controls (Appendix E, E)
197 cover all 6 frontier models.

198 **H1 ruled out as the limiting factor ($L4 \approx 1.0$).** GT-SelfStats reaches $CBF \geq 0.992$ on every one
199 of the 15 models (worst gap to perfection 0.008 on Claude-Sonnet-4-6). When the true per-domain
200 statistics are placed in the prompt, every tested model can format and emit them: the L2–L1 gap is
201 not located in the output channel.

202 **H2 splits cleanly by model family: GPT well-calibrated, Claude inflated.** *Frontier-tier means.*
203 The within-tier means in Table 2 make the picture sharp: across all 6 frontier models, LSA averages
204 CBF 0.748 vs. ZS-SK’s 0.640 and the best algorithmic L2 method’s 0.634 (UF-PDC)—a +0.108
205 margin of LSA over ZS that flips the previous 10×50 result on Claude. LSA wins on Claude-Sonnet-
206 4-6 (0.720), Claude-Opus-4-6 (0.680), Claude-Opus-4-7 (0.641), and GPT-5.5 (0.902, the highest
207 non-oracle value in the benchmark); ZS wins on GPT-5.4 (0.795) and GPT-5.4-mini (0.769), where
208 ZS estimates are well-calibrated. The 4–2 frontier split favoring LSA is driven entirely by the three
209 Claude models, where ZS estimates are systematically inflated relative to true per-domain accuracy
210 and fall *below* both NoMem and the algorithmic L2 cluster. LSA, by contrast, exceeds all algorithmic
211 L2 methods on *every* frontier row (smallest margin: +0.048 on Claude-Opus-4-7 vs UF-PDC).

212 *Strong-mid tier and individual splits.* ZS sees *no* session history, *no* feedback, yet the 50×50 update
213 places ZS first on the strong-mid tier (n=3): tier mean 0.744 (ZS) vs. 0.717 (best L2: EM-Joint).
214 Individual outcomes are split: ZS wins on Qwen-Plus (0.735 vs LSA 0.435 catastrophic failure);
215 LSA wins on Qwen-Max (0.781 vs ZS 0.766, +0.015); Qwen3.5-27B is a three-way tie at 0.731
216 (NoMem = ZS = LSA).

217 *Mid and weak tiers.* For mid-tier models the prior is partial; for the weakest models the prior matches
 218 NoMemory exactly (Mistral, GLM-5: $ZS = \text{NoMem} = \text{trivial } \hat{p} \approx 0.5$ anchor). The bottleneck
 219 location is therefore agent-dependent: for frontier and strong-mid agents the prior *is* the answer;
 220 for mid-tier agents EM-Joint and LSA tie modestly above the prior; for weak agents both prior and
 221 feedback inference are uninformative and the trivial $\hat{p} = 0.5$ NoMemory anchor wins by default.

222 **H3 localized: first-order online learning fails; second-order estimators and in-context aggregation**
 223 **succeed.** *First-order methods collapse on the frontier.* The five *first-order* L2 feedback-learning
 224 methods (UF-PDC, Platt-Online, HistBin-Online, Bayes-Domain, EM-Joint) cluster tightly in 0.69–
 225 0.72 mean CBF and fail to exceed L1 ZS on the frontier (mean L2 0.628–0.634 vs L1 ZS 0.640).

226 *Second-order estimators and in-context aggregation substantially exceed L1.* Three classes of
 227 methods *do* exceed the L1 prior on frontier ($n=6$): GLAD reaches mean CBF 0.712 (6/6 wins;
 228 Wilcoxon $p = 0.031$, binomial $p = 0.016$, Cohen’s $d = 1.42$), DS-asym 0.755 (5/6 wins; Wilcoxon
 229 one-sided $p = 0.047$, $d = 1.66$), and LSA 0.748 (mean $+0.108$, 4/6 wins, Wilcoxon $p = 0.156$
 230 inflated by GPT models where ZS is well-calibrated). GLAD’s improvement is consistent in direction
 231 across the 6 frontier models (per-model deltas $+0.028$ to $+0.118$; bootstrap 95% CI of the mean
 232 delta $[+0.045, +0.099]$ from 10K resamples; full per-model breakdown in Appendix D); with six
 233 paired observations, we read this as a direction-of-effect test rather than a population-level magnitude
 234 estimate. On non-frontier tiers, EM-Joint emerges as the leader (mid $+0.012$ over NoMem; weak
 235 $+0.006$), winning outright on Llama-3-8B (0.732) and Qwen2.5-7B (0.753).

236 *Refined H3 conclusion.* The original H3 framing (“noisy-feedback online inference is the bottleneck”)
 237 stands only for first-order single-rater estimators; second-order estimators and in-context aggregation
 238 substantially exceed L1, but none reaches the L3 ceiling—the actual bottleneck is the residual L2-
 239 vs-L3 gap of ≈ 0.20 between the strongest second-order method (DS-asym 0.755) and OraclePDC
 240 (1.000). The 50×50 protocol also reveals that several previously-best L2 baselines from the 10×50
 241 sample were within sampling noise of NoMemory; sample-size matters.

242 **Capability spectrum and best-baseline assignment.** With the protocol (Table 2), the “best” non-
 243 oracle baseline per row (column “Best” in Table 2, defined as the argmax across the 8 baseline columns
 244 L0 NoMem + L1 ZS + $6 \times$ L2) splits three ways across 15 models: LSA wins on 6 (DeepSeek-V3.2,
 245 Qwen-Max, GPT-5.5, Claude-Sonnet-4-6, Claude-Opus-4-6, Claude-Opus-4-7), EM-Joint wins on 5
 246 (Mistral-7B, Llama-3-8B, GLM-5, Qwen2.5-7B, Qwen3.5-27B—note that on Mistral-7B the EM-
 247 NoMem-ZS-LSA spread is within ± 0.003 , so the EM win is essentially a tie with the trivial $\hat{p} \approx 0.5$
 248 anchor), and ZS-SK wins on 4 (Qwen-Turbo, Qwen-Plus, GPT-5.4-mini, GPT-5.4). NoMemory does
 249 not strictly win on any model under this argmax. *Note:* the “feedback as friend or foe” analysis in
 250 Appendix E uses the orthogonal head-to-head ZS-vs-LSA criterion; a model can therefore appear
 251 under “LSA-dominant best-of-8” here yet “ZS-wins-over-LSA” in the head-to-head, since the criteria
 252 differ.

253 4.3 Diagnosis via the Ladder Visualization

254 Figure 2 compresses the per-row Table 2 story into vertical structure that reads each adjacent gap as a
 255 sub-capability test: the bottleneck is not the output channel ($L4 \approx L3 \approx 1.0$) but noise-robust online
 256 inference ($L2 \rightarrow L3$ gap ≈ 0.30 across all 15 models).

257 Reading the ladder by adjacent comparisons: $L4$ vs $L3$ ($< 1\%$ gap)—the output channel is not the
 258 constraint, ruling out H1. $L3$ vs $L2$ (~ 0.30 gap)—the empirical bottleneck: with clean per-turn
 259 labels agents reach 1.0, but no method fully recovers the signal from noisy LLM-decoded feedback,
 260 confirming H3. $L2$ vs $L1$ ($n=15$: 0.674–0.686 vs 0.685; $n=6$ frontier: 0.628–0.634 vs 0.640)—L1
 261 ZS ties or exceeds the algorithmic L2 cluster on average; only second-order estimators (DS-asym,
 262 GLAD) and in-context LSA exceed L1 on the frontier (Appendix E, E).

263 **A diagnostic per-model selection rule (R4).** The H2 family-split above motivates a simple sanity-
 264 check rule rather than a new estimator: **R4** (per-model best-of-two between the prior and in-context
 265 aggregation) selects L1 (ZS-SK) if $ZS(m) > \text{NoMem}(m)$, otherwise L2 (LSA). The “R4” label is
 266 internal to our four-rule analysis pipeline (R1–R4 documented in Appendix O); we report only R4
 267 here because the R1–R3 rules are subsumed. R4 attains mean CBF 0.722 across $n=15$ and closes
 268 $\approx 72\%$ of the always-ZS-to-per-model-oracle gap, but its per-model sign tests against always-ZS




Figure 2: **The L0–L4 ladder localizes the bottleneck.** Mean CBF per baseline level on (a) all 15 models and (b) the 6 frontier models. **(a)** L2→L3 gap ≈ 0.30 (best L2 LSA 0.702 vs. OraclePDC 1.000): first-order L2 methods do not recover the signal from noisy feedback. **(b)** On the frontier, the first-order L2 cluster (Platt/HistBin/Bayes/EM/UF) collapses to NoMem (0.628–0.634); only second-order estimators (DS-Asym 0.755, GLAD 0.712) and in-context LSA (0.748) exceed the L1 ZS prior (0.640), but all fall well short of OraclePDC. Dotted vertical line marks the L0 NoMemory anchor.

($p=0.22$) and always-LSA ($p=0.69$) are non-significant: the rule *redistributes* wins, it does not dominate. We present R4 as a benchmark-validation lower bound—any new method should clear it—not as a contribution; the empirical frontier ceiling remains DS-Asym at CBF 0.755. Full LOOCV, per-model breakdown, and sign-test arithmetic in Appendix O.

4.4 Cross-Task Heterogeneity: Per-Task Boundaries Are Structurally Necessary

Across MMLU, GSM8K, HumanEval, and BFCL the same model exhibits qualitatively different calibration regimes. Within the protocol-controlled extension suite (GSM8K/HumanEval/BFCL under the same single-shot two-step protocol), 12/15 models show Gap-sign reversals across tasks, and a Platt scaling learned on MMLU *harms* GSM8K for the majority of models. Per-task boundary modeling is therefore a structural necessity, not a convenience. Full per-task table, controlled-protocol subset analysis, and Platt-transfer details: Appendix Q.1 and Table 6.

5 Limitations and Future Work

Scope of evaluation. SelfBound evaluates per-domain self-knowledge under a 50-turn MMLU protocol with a controlled synthetic feedback channel (75% in-expertise, 46% out, cross-family routing to mitigate dual-role bias). The three extension tasks (GSM8K, HumanEval, BFCL) use a single-shot two-step protocol because they lack the multi-domain structure CBF requires; the cross-task heterogeneity finding in §4.4 is therefore independently validated via a protocol-controlled subset analysis (12/15 sign-flip rate on the three same-protocol tasks; Appendix Q.1).

Validity bounds. L1 ZS-SK measures per-domain *capability estimation* rather than self-knowledge in the strict sense; MMLU-Pro and BBH contamination controls (Appendix E) attenuate but do not eliminate benchmark-name memorization. The frontier sub-sample of $n=6$ supports direction-of-effect reporting (sign consistency, bootstrap CIs of paired deltas; Appendix D) rather than population-level magnitude. The headline conclusions (H1 does not bind, H2 family-splits, H3 is second-order tractable) do not depend on the literal reading of either.

Next steps.

- *Human-feedback pilot:* replace the synthetic channel with real users and check whether the H3 ranking transfers.
- *Full ladder on extension tasks:* replicate L0–L4 on GSM8K/HumanEval/BFCL.

- 297 • *Frontier-slice extension*: scale to $n \approx 18$ to refine the DS-asym vs. GLAD comparison (per
298 Appendix D power calculation).
- 299 • *Open-ended multi-step agents*: extend to SWE-Bench-style settings.

300 6 Conclusion

301 We asked where, mechanistically, LLM-agent self-modeling fails when feedback is noisy. The
302 five-level oracle ladder of SelfBound rules out the output channel (H1: $\text{CBF} \geq 0.99$ given ground-
303 truth prompts), shows the zero-shot prior to be partial and family-split rather than the bottleneck
304 (H2), and localizes the failure to noise-robust second-order online inference (H3): first-order single-
305 rater learners cannot recover per-domain truth from noisy feedback, while explicit second-order
306 estimators—Dawid–Skene-asym and GLAD—substantially close the L1→L3 gap (GLAD beats
307 L1 on all six frontier models, $p = 0.031$) without yet eliminating it. We name this failure mode
308 **Self-Boundary Collapse** and release the ladder, the CBF metric, the cross-family user simulator,
309 and an open leaderboard so that the community can evaluate new second-order estimators against a
310 precise, oracle-bracketed target rather than ad-hoc baselines. The R4 parameter-free selection rule
311 (Appendix O) serves as the validation lower bound any submitted method should clear. The open
312 problem is concrete: a noise-robust online estimator that closes the remaining L2→L3 gap—roughly
313 0.30 CBF on the frontier—would constitute a measurable advance in agent self-modeling.

314 References

- 315 [1] Christopher Ackerman. Evidence for limited metacognition in llms. *Proceedings of ICLR*,
316 2026.
- 317 [2] Gustaf Ahdritz, Tian Qin, Nikhil Vyas, Boaz Barak, and Benjamin L. Edelman. Distinguishing
318 the knowable from the unknowable with language models. *Proceedings of ICML*, 2024.
- 319 [3] Casey O. Barkan, Sid Black, and Oliver Sourbut. Do large language models know what they
320 are capable of? *Proceedings of ICLR*, 2026.
- 321 [4] Felix J Binder, James Chua, Tomek Korbak, Henry Sleight, John Hughes, Robert Long, Ethan
322 Perez, Miles Turpin, and Owain Evans. Looking inward: Language models can learn about
323 themselves by introspection. *Proceedings of ICLR*, 2025.
- 324 [5] Jiawei Chen, Ruoxi Xu, Boxi Cao, Ruotong Pan, Yunfei Zhang, Yifei Hu, Yong Du, Tingting
325 Gao, Yaojie Lu, Yingfei Sun, Xianpei Han, Le Sun, Xiangyu Wu, and Hongyu Lin. Towards
326 real-world human behavior simulation: Benchmarking large language models on long-horizon,
327 cross-scenario, heterogeneous behavior traces. *arXiv preprint arXiv:2604.08362*, 2025.
- 328 [6] Lihu Chen, Alexandre Perez-Lebel, Fabian M Suchanek, and Gaël Varoquaux. Reconfidencing
329 llms from the grouping loss perspective. *Findings of EMNLP*, 2024.
- 330 [7] Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, et al. Evaluating large language models
331 trained on code. *arXiv preprint arXiv:2107.03374*, 2021.
- 332 [8] Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser,
333 Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, et al. Training verifiers to
334 solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.
- 335 [9] Shizhe Diao et al. R-tuning: Instructing large language models to say i don’t know. *Proceedings*
336 *of NAACL (Outstanding Paper)*, 2024.
- 337 [10] Y. Dong, X. Jiang, et al. RL-plus: Countering capability boundary collapse of llms in reinforce-
338 ment learning. *arXiv preprint arXiv:2508.00222*, 2025.
- 339 [11] Ke Fang, Tianyi Zhao, and Lu Cheng. Credence calibration game? calibrating large language
340 models through structured play. *arXiv preprint arXiv:2508.14390*, 2025.
- 341 [12] Jiahui Geng et al. A survey of confidence estimation and calibration in large language models.
342 2024.

- [13] Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q Weinberger. On calibration of modern neural networks. In *ICML*, 2017.
- [14] Haixia Han, Tingyun Li, Shisong Chen, Jie Shi, Chengyu Du, Yanghua Xiao, Jiaqing Liang, and Xin Lin. Enhancing confidence expression in large language models through learning from past experience. *arXiv preprint arXiv:2404.10315*, 2024.
- [15] Dan Hendrycks et al. Measuring massive multitask language understanding. In *ICLR*, 2021.
- [16] Yuanzhe Hu, Yu Wang, and Julian McAuley. Evaluating memory in llm agents via incremental multi-turn interactions. *Proceedings of ICLR*, 2026.
- [17] Li Ji-An, Hua-Dong Xiong, Robert C Wilson, Marcelo G Mattar, and Marcus K Benna. Language models are capable of metacognitive monitoring and control of their internal activations. *Proceedings of NeurIPS*, 2025.
- [18] Saurav Kadavath et al. Language models (mostly) know what they know. *arXiv preprint arXiv:2207.05221*, 2022.
- [19] Sanyam Kapoor, Nate Gruver, Manley Roberts, Katherine Collins, Arka Pal, Umang Bhatt, Adrian Weller, Samuel Dooley, Micah Goldblum, and Andrew Gordon Wilson. Large language models must be taught to know what they don’t know. *Proceedings of NeurIPS*, 2024.
- [20] Polina Kirichenko, Mark Ibrahim, Kamalika Chaudhuri, and Samuel J. Bell. Abstentionbench: Do reasoning llms know what they don’t know? *Proceedings of NeurIPS*, 2025.
- [21] Rudolf Kulhavý and Michael B Zarrop. On a general concept of forgetting. *International Journal of Control*, 58(4):905–924, 1993.
- [22] Rudolf Laine, Bilal Chughtai, Jan Betley, Kaivalya Hariharan, Jeremy Scheurer, Mikita Balesni, Marius Hobbhahn, Alexander Meinke, and Owain Evans. Towards a situational awareness benchmark for llms. *arXiv preprint arXiv:2407.04694*, 2024.
- [23] Jixuan Leng, Chengsong Huang, Banghua Zhu, and Jiabin Huang. Taming overconfidence in llms: Reward calibration in rlhf. *Proceedings of ICLR*, 2025.
- [24] Moxin Li, Yong Zhang, Jingxuan Wen, et al. Knowledge boundary of large language models: A survey. *Proceedings of ACL*, 2025.
- [25] Siheng Li et al. Know your limits: A survey of abstention in large language models. *Transactions of the Association for Computational Linguistics*, 2025.
- [26] Zhuoyi Li et al. Memos: An operating system for memory-augmented generation in large language models. *arXiv preprint arXiv:2507.03724*, 2025.
- [27] Percy Liang, Rishi Bommasani, Tony Lee, Dimitris Tsipras, Dilara Soylu, Michihiro Yasunaga, Yian Zhang, Deepak Narayanan, Yuhuai Wu, Ananya Kumar, et al. Holistic evaluation of language models. *arXiv preprint arXiv:2211.09110*, 2022.
- [28] Stephanie Lin, Jacob Hilton, and Owain Evans. Truthfulqa: Measuring how models mimic human falsehoods. *Proceedings of ACL*, 2022.
- [29] Jiayin Liu, Ruochen Zhang, and Yejin Choi. Implicit user feedback in human-llm dialogues: Informative to understand users yet noisy as a learning signal. *Proceedings of EMNLP*, 2025.
- [30] Tennison Liu and Mihaela van der Schaar. Position: Truly self-improving agents require intrinsic metacognitive learning. *Proceedings of ICML*, 2025.
- [31] Chang Ma, Junlei Zhang, Zhihao Zhu, Cheng Yang, Yujiu Yang, Yaohui Jin, Zhenzhong Lan, Lingpeng Kong, and Junxian He. Agentboard: An analytical evaluation board of multi-turn llm agents. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.

- [32] Putra Manggala, Atalanti Mastakouri, Elke Kirschbaum, Shiva Prasad Kasiviswanathan, and Aaditya Ramdas. Qa-calibration of language model confidence scores. *Proceedings of ICLR*, 2025.
- [33] Mem0 Team. Mem0: Building production-ready ai agents with scalable long-term memory. *arXiv preprint arXiv:2504.19413*, 2025.
- [34] Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to follow instructions with human feedback. *Advances in Neural Information Processing Systems*, 35:27730–27744, 2022.
- [35] Charles Packer, Sarah Wooders, Kevin Lin, et al. Memgpt: Towards llms as operating systems. *arXiv preprint arXiv:2310.08560*, 2023.
- [36] John Platt. Probabilistic outputs for support vector machines and comparisons to regularized likelihood methods. In *Advances in Large Margin Classifiers*. MIT Press, 1999.
- [37] Rafael Rafailov, Archit Sharma, Eric Mitchell, Christopher D Manning, Stefano Ermon, and Chelsea Finn. Direct preference optimization: Your language model is secretly a reward model. *Advances in Neural Information Processing Systems*, 36, 2024.
- [38] Vishnu Sarukkai, Zhiqiang Xie, and Kayvon Fatahalian. Self-generated in-context examples improve llm agents for sequential decision-making tasks. *Advances in Neural Information Processing Systems*, 2025.
- [39] Maohao Shen, Subhro Das, Kristjan Greenewald, Prasanna Sattigeri, Gregory Wornell, and Soumya Ghosh. Thermometer: Towards universal calibration for large language models. *Proceedings of ICML*, 2024.
- [40] Mirac Suzgun, Nathan Scales, Nathanael Sch"arli, Sebastian Gehrmann, Yi Tay, Hyung Won Chung, Aakanksha Chowdhery, Quoc V Le, Ed H Chi, Denny Zhou, Hanie Sedghi, Hessam Bagherinezhad, Yasaman Bahri, Dale Schuurmans, Olivier Bousquet, Quoc V Le, and Jason Wei. Challenging big-bench tasks and whether chain-of-thought can solve them. *arXiv preprint arXiv:2210.09261*, 2022.
- [41] Juntao Tan et al. In prospect and retrospect: Reflective memory management for long-term personalized dialogue agents. *Proceedings of ACL*, 2025.
- [42] Benedikt Traub et al. Overcoming common flaws in the evaluation of selective classification systems. *Proceedings of NeurIPS*, 2024.
- [43] Aaron David Tucker, Kianté Brantley, Adam Cahall, and Thorsten Joachims. Coactive learning for large language models using implicit user feedback. *Proceedings of ICML*, 2024.
- [44] Jason Wei et al. Simpleqa: Measuring short-form factuality in large language models. *OpenAI Technical Report*, 2024.
- [45] Miao Xiong, Zhiyuan Hu, Xinyang Lu, Yifei Li, Jie Fu, Junxian He, and Bryan Hooi. Can llms express their uncertainty? an empirical evaluation of confidence elicitation in llms. In *ICLR*, 2024.
- [46] Tianyang Xu, Shujin Wu, Shizhe Diao, Xiaoze Liu, Xingyao Wang, Yangyi Chen, and Jing Gao. Saysself: Teaching llms to express confidence with self-reflective rationales. *Proceedings of EMNLP*, 2024.
- [47] Wujiang Xu, Zujie Liang, Kai Mei, Hang Gao, Juntao Tan, and Yongfeng Zhang. A-mem: Agentic memory for llm agents. *Proceedings of NeurIPS*, 2025.
- [48] Fanjia Yan, Huanzhi Mao, Charlie Cheng-Jie Ji, Tianjun Zhang, Shishir G. Patil, Ion Stoica, and Joseph E. Gonzalez. Berkeley function calling leaderboard. <https://gorilla.cs.berkeley.edu/leaderboard.html>, 2024.

- [49] Shunyu Yao, Noah Shinn, Pedram Razavi, and Karthik Narasimhan. *tau-bench: A benchmark for tool-agent-user interaction in real-world domains. arXiv preprint arXiv:2406.12045*, 2024.
- [50] Xunjian Yin, Xu Zhang, Jie Ruan, and Xiaojun Wan. Benchmarking knowledge boundary for large language models: A different perspective on model evaluation. *Proceedings of ACL*, 2024.
- [51] Zhangyue Yin, Qiushi Sun, Qipeng Guo, Jiawen Wu, Xipeng Qiu, and Xuanjing Huang. Do large language models know what they don’t know? In *Findings of ACL*, 2023.
- [52] Bianca Zadrozny and Charles Elkan. Obtaining calibrated probability estimates from decision trees and naive bayesian classifiers. In *Proceedings of ICML*, 2001.
- [53] Xuhui Zhou, Weiwei Sun, Qianou Ma, et al. Mind the sim2real gap in user simulation for agentic tasks. *arXiv preprint*, 2026.
- [54] Zijian Zhou, Ao Qu, Zhaoxuan Wu, et al. Mem1: Learning to synergize memory and reasoning for efficient long-horizon agents. *Proceedings of ICLR*, 2026.

NeurIPS Paper Checklist

1. **Claims**
Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?
Answer: [Yes]
Justification: Abstract and §1 state the H1/H2/H3 decomposition, the L0–L4 ladder, the CBF metric, and the empirical findings (H1 does not bind, H2 family-splits, H3 localized to second-order inference). All claims are supported by the experiments in §4.2–§4.3 and the appendices.
2. **Limitations**
Question: Does the paper discuss the limitations of the work performed by the authors?
Answer: [Yes]
Justification: §5 discusses scope, n=6 frontier sample, validity of L1 channel, simulated-user limits, and pending human-feedback validation.
3. **Theory Assumptions and Proofs**
Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?
Answer: [N/A]
Justification: The paper is empirical/diagnostic; it does not introduce new theoretical results.
4. **Experimental Result Reproducibility**
Question: Does the paper fully disclose all the information needed to reproduce the main experimental results?
Answer: [Yes]
Justification: Full protocol (50×50), persona parameterisation, baseline algorithms, and decoder configuration are in §3 and Appendix B, C. Open-weights models (Mistral-7B, Llama-3-8B) anchor reproducibility via pinned HF commits.
5. **Open Access to Data and Code**
Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results?
Answer: [Yes]
Justification: Code, 47K-turn JSONL traces, per-model domain profiles, and Croissant metadata released in the anonymized repository (URL in OpenReview submission fields).
6. **Experimental Setting/Details**
Question: Does the paper specify all the training and test details necessary to understand the results?
Answer: [Yes]
Justification: Setup, model versions, API access dates, decoder, parser specification, and per-baseline hyperparameters are documented in §4.1, Appendix B, C, S.

- 484 **7. Experiment Statistical Significance**
 485 **Question:** Does the paper report error bars suitably and correctly?
 486 **Answer:** [Yes]
 487 **Justification:** Bootstrap 95% CIs, Wilcoxon and binomial tests, Cohen’s d , and per-model
 488 bootstrap stability are reported in §4.2, Appendix D, G. The $n=6$ frontier statistics are
 489 explicitly framed as direction-of-effect tests.
- 490 **8. Experiments Compute Resources**
 491 **Question:** For each experiment, does the paper provide sufficient information on the
 492 computer resources?
 493 **Answer:** [Yes]
 494 **Justification:** Open-weights inference uses vLLM on a single A6000 (48GB); API calls are
 495 documented in Appendix S with access dates. Total compute: $\approx 47\text{K}$ LLM calls across 15
 496 models $\times$ 10 baselines.
- 497 **9. Code of Ethics**
 498 **Question:** Does the research conducted in the paper conform, in every respect, with the
 499 NeurIPS Code of Ethics?
 500 **Answer:** [Yes]
 501 **Justification:** No human subjects; no personally identifiable data; uses public benchmarks
 502 (MMLU, GSM8K, HumanEval, BFCL) under their respective licenses with attribution.
- 503 **10. Broader Impacts**
 504 **Question:** Does the paper discuss both potential positive societal impacts and negative
 505 societal impacts of the work performed?
 506 **Answer:** [Yes]
 507 **Justification:** Improving LLM agent self-knowledge supports safer deployment in long-
 508 horizon settings (positive); however, advances in self-modeling could also be exploited to
 509 mask uncertainty in adversarial contexts (§5).
- 510 **11. Safeguards**
 511 **Question:** Does the paper describe safeguards that have been put in place for responsible
 512 release of data?
 513 **Answer:** [N/A]
 514 **Justification:** Released artifacts contain only LLM outputs on public benchmarks; no
 515 high-risk content.
- 516 **12. Licenses for Existing Assets**
 517 **Question:** Are the creators or original owners of assets used in the paper, properly credited
 518 and are the license and terms of use explicitly mentioned and respected?
 519 **Answer:** [Yes]
 520 **Justification:** MMLU (MIT), GSM8K (MIT), HumanEval (MIT), BFCL (Apache-2.0)
 521 attributed in §3 and the data hosting documentation.
- 522 **13. New Assets**
 523 **Question:** Are new assets introduced in the paper well documented and is the documentation
 524 provided alongside the assets?
 525 **Answer:** [Yes]
 526 **Justification:** Code released under MIT, dataset under CC-BY-4.0; Croissant metadata,
 527 README, MANIFEST, and per-baseline algorithm documentation included.
- 528 **14. Crowdsourcing and Research with Human Subjects**
 529 **Question:** For crowdsourcing experiments and research with human subjects, does the paper
 530 include the full text of instructions given to participants?
 531 **Answer:** [N/A]
 532 **Justification:** No human subjects; the simulated user is an LLM persona (full prompt
 533 template in Appendix B).
- 534 **15. Institutional Review Board (IRB) Approvals or Equivalent**
 535 **Question:** Does the paper describe potential risks incurred by study participants, and obtain
 536 appropriate IRB approvals?
 537 **Answer:** [N/A]
 538 **Justification:** No human subjects.

539 **16. Declaration of LLM usage**
540 **Question:** Does the paper describe the usage of LLMs?
541 **Answer:** [\[Yes\]](#)
542 **Justification:** LLMs are the subject of the benchmark and serve as agents, simulated users,
543 and feedback decoders; full configuration documented in §3, Appendix S.

544 Datasheet for Datasets

545 We follow the structure proposed by Gebru et al. (2018).

546 Motivation

547 SelfBound was created to fill a measurement gap in agentic LLM evaluation: existing benchmarks
548 (MMLU, BIG-Bench, AgentBench) score one-shot accuracy but say nothing about whether an agent
549 can *learn its own per-domain capability boundary* across repeated interactions.

550 Composition

551 20 MMLU subjects $\times$ 25 questions = 500 main pool, supplemented by GSM8K, HumanEval, BFCL
552 Simple v3. 50-turn dialogue sessions, 10–50 sessions per (agent, persona). Each instance is a JSONL
553 session log including question, agent answer, agent confidence, persona reaction, ground-truth
554 grading.

555 Collection Process

556 Questions from public test splits of MMLU (Hendrycks 2021), GSM8K (Cobbe 2021), HumanEval
557 (Chen 2021), BFCL Simple v3 (gorilla-llm). Personas synthesized via cross-family LLM (Qwen-
558 Max for Qwen agents, GPT-5.4-mini for Claude agents, Claude-Sonnet for GPT agents) to control
559 evaluator-bias.

560 Preprocessing

561 Per-model tier stratification: held-out validation accuracy on disjoint 10-question/subject split, then
562 Strong/Mid/Weak partitioning of the 20 subjects. This per-model stratification (rather than a global
563 tier) makes boundary-learning non-trivial.

564 Uses

565 Intended: (a) self-knowledge memory evaluation; (b) noise-robust online inference algorithms; (c)
566 longitudinal frontier-model calibration tracking. Unsuitable: ranking models on raw accuracy.

567 Recommended Splits and Stratification

568 The 20 MMLU subjects are partitioned per-model into Strong / Mid / Weak tiers using each model’s
569 held-out validation accuracy on a disjoint 10-question/subject calibration split. We do *not* provide a
570 fixed train/test split: every reported number uses the same per-model stratification protocol (described
571 in Appendix B) so that downstream researchers can reproduce or substitute their own model. The
572 four cross-task evaluations (GSM8K, HumanEval, BFCL, MMLU-Pro) use the upstream provider’s
573 official test split.

574 Errors, Noise, and Known Issues

575 Three classes of known issues should be considered when interpreting our reported numbers:

- 576 • *Algorithmic-baseline implementation note* (§A): all algorithmic L2 baselines compute their
577 own per-domain posteriors; we document the implementation in Appendix A so that re-
578 implementations can verify they reproduce the same behavior.
- 579 • *GSM8K Opus-4-7 truncation*: 158 of 300 GSM8K questions were truncated by the original
580 max_tokens limit on Claude-Opus-4-7; the reported pass-rate is on the 142-question answered
581 subset (see Table 6).

582 We surface these openly because their detection and disclosure is itself a methodological contribution
583 of this benchmark; future versions will document any newly identified issues here.

584 **Maintenance Contact**

585 Maintainer email and a dedicated Slack/Discord workspace will be made available upon de-
586 anonymization. For double-blind review, please open an issue at the anonymised repository above;
587 we monitor it during the rebuttal window.

588 **Distribution and Maintenance**

589 Released artifacts under a dual SPDX license: code under MIT, data under CC-BY-4.0. Specifically:
590 protocol code, baseline implementations, evaluation scripts (MIT); raw session JSONL (~47K
591 interactions including all per-turn agent outputs, simulated user reactions, and ground-truth grading),
592 per-model domain profiles, and the 20-MMLU-subject test pool (CC-BY-4.0 with attribution to
593 upstream MMLU/GSM8K/HumanEval/BFCL providers). 6-month review cycle; new model additions
594 via PR (see Submission Protocol).

595 **Maintenance Plan**

596 **Versioning**

597 Semantic over the protocol: v3 (50×50) is current; v2 (10×50) supported as lite. Question pool
598 versioning is independent.

599 **Leaderboard Hosting**

600 A HuggingFace Spaces leaderboard (URL released on acceptance). CBF and ancillary metrics (ECE,
601 Brier, AUROC, HR) auto-recomputed from submitted JSONL via the public scorer script in the repo.

602 **Submission Protocol**

603 PR-based. Each submission must include:

- 604 1. **Full session JSONL files** (one per agent-persona pair) following the schema:
- ```
605 {
606 "session_id": str, "agent_model": str, "persona": dict,
607 "log": [
608 {"turn": int, "domain": str, "tier": "strong|mid|weak",
609 "question": str, "agent_answer": str, "agent_confidence": float,
610 "user_response": str, "gt_correct": bool}
611]
612 }
```
- 613 2. Reference implementation of the agent’s online-inference component (Python module exposing  
614 `predict_confidence(turn)` and `self_assess(domain)`).
- 615 3. Paper or technical report reference.

616 Aggregate-scores-only submissions are not merged; we recompute all metrics from raw JSONL via  
617 the public scorer.

618 **Deprecation Policy for API Models**

619 Deprecated endpoints’ rows move to *archived* table with frozen score and deprecation date; not  
620 deleted. Preserves longitudinal comparability.

621 **Contact**

622 Anonymized for review; withheld for double-blind review.

## 623 **Ethics Statement**

### 624 **Simulated Users**

625 All "users" are LLM-simulated personas, not human participants. No PII. Programmatically-generated  
626 expertise labels. Does not require IRB review under our institution's policy.

### 627 **LLM-as-Evaluator and Cross-Family Routing**

628 Dual-role bias mitigation via cross-family persona-generation (forbidden: Qwen-Max evaluates  
629 Qwen agents). Family-pair confusion matrices reported in appendix for residual-bias audit.

### 630 **Provider Terms of Service**

631 LLM API outputs used strictly under research-evaluation terms. Released JSONL not redistributed  
632 for model training; license header explicit.

### 633 **No Deceptive Prompting**

634 All "role-playing" instructions stay within LLM-to-LLM simulation envelope.

### 635 **Potential Misuse**

636 Most plausible misuse falls in two categories: (a) over-interpretation of headline CBF as a general  
637 "self-awareness" ranking suitable for safety certification (CBF measures *per-domain accuracy es-*  
638 *timation*, not *overall trustworthiness* — using it to certify deployment in clinical / legal / financial  
639 high-stakes contexts would be a category error); (b) deployment misuse where a model with high  
640 CBF is presumed safe to operate without human oversight in domains beyond the 20 MMLU subjects  
641 evaluated. Addressed by framing as diagnostic instrument; per-tier breakdowns discourage scalar  
642 comparisons.

## 643 **Broader Impact**

### 644 **Positive Impact**

645 Enables systematic evaluation of agent self-knowledge—prerequisite for safer customer-service  
646 / financial-advisory / clinical-triage deployment where over-confident answers carry asymmetric  
647 downside cost.

### 648 **Reproducibility Cost**

649 Full v3 50×50 pipeline: ~USD 200–400 at 2026 rates across 13–15 frontier models. Lite v2 10×50  
650 reproduces headline ordering at ~USD 40. Raw JSONL released for re-scoring without re-running.

### 651 **Evaluator-Model Bias**

652 Cross-family policy mitigates but does not eliminate dual-role bias. Residual vendor-cluster effect  
653 possible. High-stakes comparisons should cross-check under  $\geq 2$  persona-generator families.

### 654 **Benchmark Overfitting**

655 Goodhart-style risk: methods optimizing for CBF rather than the underlying construct. Mitigated by  
656 held-out 5-subject transfer split; large dev-vs-transfer gaps flagged on leaderboard.

### 657 **Longitudinal Stability**

658 API-only frontier rows: ~12-month half-life. Open-weights anchors (Mistral-7B, Llama-3-8B)  
659 permanently included for fixed-reference recalibration.

**Extended snapshot manifest.** Table S lists access dates only; it does not establish whether each model is content-addressable. Table 3 fills this gap by recording, for every one of the 15 models, the exact identifier sent to the endpoint, the version-pin scheme observable from the response payload, and a per-model reproducibility risk classification.

Table 3: Extended model snapshot manifest. Repro Risk reflects whether a third party in 2027 could obtain the same weights from the listed identifier. LOW = HuggingFace commit hash anchors the weights; MED = a content-addressable anchor exists but the served endpoint is a floating alias; HIGH = the endpoint returns no version-suffixed identifier and no public open-weights anchor exists.

| Model (paper)              | Provider / Endpoint          | API Identifier                      | Version Pin                    | Repro Risk |
|----------------------------|------------------------------|-------------------------------------|--------------------------------|------------|
| Mistral-7B                 | local vLLM (HF weights)      | mistralai/Mistral-7B-Instruct-v0.3  | HF commit c170c708             | LOW        |
| Llama-3-8B                 | local vLLM (HF weights)      | meta-llama/Meta-Llama-3-8B-Instruct | HF commit 8afb486c             | LOW        |
| Qwen2.5-7B                 | DashScope                    | qwen2.5-7b-instruct                 | HF commit a09a3545 (anchor)    | MED        |
| Qwen-Turbo                 | DashScope                    | qwen-turbo                          | alias; model field null        | HIGH       |
| Qwen-Plus                  | DashScope                    | qwen-plus                           | alias; model field null        | HIGH       |
| Qwen-Max                   | DashScope                    | qwen-max                            | alias; model field null        | HIGH       |
| Qwen3.5-27B                | DashScope                    | qwen3.5-27b                         | alias; no public HF release    | HIGH       |
| DeepSeek-V3.2 <sup>§</sup> | DashScope                    | deepseek-v3 (sent)                  | alias; provider-rehosted ckpt  | HIGH       |
| GLM-5                      | DashScope                    | glm-5                               | alias; provider-rehosted ckpt  | HIGH       |
| GPT-5.4-mini               | proxy zhongzhuan / responses | gpt-5.4-mini                        | echoes gpt-5.4-mini-2026-03-17 | MED        |
| GPT-5.4                    | proxy zhongzhuan / responses | gpt-5.4                             | alias; no date suffix returned | HIGH       |
| GPT-5.5                    | proxy zhongzhuan / responses | gpt-5.5                             | alias; no date suffix returned | HIGH       |
| Claude-Sonnet-4-6          | Anthropic-compatible proxy   | claude-sonnet-4-6                   | alias; no date suffix returned | HIGH       |
| Claude-Opus-4-6            | Anthropic-compatible proxy   | claude-opus-4-6                     | alias; no date suffix returned | HIGH       |
| Claude-Opus-4-7            | proxy zhongzhuan / messages  | claude-opus-4-7                     | alias; no date suffix returned | HIGH       |

**API-versioning gaps and a paper-code label mismatch.** Three entries deserve foregrounding. First, an internal label mismatch: the paper refers to “DeepSeek-V3.2” (denoted <sup>§</sup>), but the string our scripts transmit to DashScope is deepseek-v3. DashScope re-hosts a checkpoint they label deepseek-v3; we lack provider documentation establishing whether this resolves to V3.0, V3.1, or V3.2-Exp. We retain the paper label as a working estimate but flag this as an unresolved provenance question. Second, 12/15 endpoints do not return version-suffixed model strings; the single exception is gpt-5.4-mini, which is auto-pinned by the proxy. Third, the three open-weights anchors are the only entries for which we can hand a future replicator a 40-character identifier sufficient to reconstruct the exact weights. Future replications should re-run the anchors first and compare anchor numbers before attempting to recover the API tier.

## A Implementation Note: Algorithmic-Baseline self\_assess Patch

**Patch description.** The five algorithmic L2 baselines (Platt-Online, HistBin-Online, Bayes-Domain, EM-Joint, UF-PDC) share a common self\_assess(domain) interface. The implementation we report uses each algorithm’s own per-domain posterior: Platt’s logistic re-calibration, Bayes-Domain’s Beta posterior mean  $\alpha/(\alpha+\beta)$ , EM-Joint’s iterated  $\hat{p}^*(d)$ , etc. We document this interface here so that downstream re-implementations reproduce the same behavior; inadvertently routing self-assessment through an LLM verbal-estimate helper produces a degenerate “algorithmic-L2-collapses-to-NoMem” pattern that is easy to mis-interpret as a property of the baseline rather than the implementation.

**Patch.** We replaced each algorithmic baseline’s self\_assess with the appropriate posterior:

- Platt-Online / HistBin-Online:  $\hat{p}(d) = \text{dom\_correct}[d]/\text{dom\_total}[d]$  (or 0.5 if no observations).
- Bayes-Domain:  $\hat{p}(d) = \alpha[d]/(\alpha[d] + \beta[d])$ .
- EM-Joint:  $\hat{p}(d) = \text{clip}(p^*[d], 10^{-3}, 1-10^{-3})$  from EM iterate.
- UF-PDC:  $\hat{p}(d) = \text{correct}[d]/\text{total}[d]$  from decoded feedback.

We also bumped the LLM-call max\_tokens from 20 to 200 for the legitimately verbal baselines (LSA, ZS-SK, GT-SelfStats) to accommodate reasoning models; ZS-SK and LSA were unaffected by the algorithmic-baseline patch but remain LLM-dependent.

**Per-baseline summary.** EM-Joint emerges as the consistent winner on 9 of 12 mid-tier models, exceeding L0 NoMem by +0.003 to +0.022. Bayes-Domain wins on GPT-5.4-mini (+0.022 over

NoMem); Platt-Online wins on GPT-5.4 and GPT-5.5 (+0.016–+0.019 over NoMem). The first-order algorithmic L2 methods cluster within  $\pm 0.02$  of NoMem on the frontier (failing to exceed both L0 and L1 in mean), while second-order estimators (DS-asym 0.755, GLAD 0.712) and in-context aggregation (LSA 0.748) cleanly exceed L1 ZS (0.640). The L2-vs-L3 gap of  $\approx 0.25$  to OraclePDC (CBF 1.000) remains the empirical bottleneck the benchmark isolates.

## B Benchmark Details: Task, Protocol, Question Pool, Personas, Metrics




Figure 3: **Capability spectrum across 15 models.** (a) MMLU accuracy spans 42–94%; family colors are consistent throughout the paper. (b) Best non-oracle CBF per model with the winning baseline label (EM/LSA/ZS). The dotted line marks the L0 NoMemory mean (n=15, 0.671). The two newly-added frontier models (GPT-5.5, Claude-Opus-4-7) anchor the high-accuracy end. Note that Claude-Opus-4-7 has lower MMLU accuracy (69%) but is included in the n=6 frontier subset for the second-order analysis.

**Task formalisation.** Given a stream  $(q_t, a_t, c_t, f_t)_{t=1}^T$  where each question  $q_t$  is drawn from a hidden domain  $d_t \in \mathcal{D}$ , the agent must construct  $\hat{p}: \mathcal{D} \rightarrow [0, 1]$  approximating the true per-domain accuracy  $p^*(d)$ . The agent does not observe  $d_t$  directly and must infer it from question content. User feedback  $f_t$  is a free-form natural-language utterance that may confirm, correct, or provide ambiguous/misleading responses. The agent outputs  $a_t$  together with confidence  $c_t \in [0, 1]$ . The challenge is that feedback reliability is domain-dependent: accurate within the user’s expertise set  $\mathcal{E} \subset \mathcal{D}$ , noisy outside it.

**Session protocol.** Each session has 50 turns. At each turn  $t$ : (1) sample  $q_t$  from the question pool with  $d_t$  hidden from the agent; (2) the agent emits  $(a_t, c_t)$ ; (3) the simulated user produces  $f_t$  conditioned on the ground-truth answer, the user’s expertise level on  $d_t$ , and a persona-consistent style; (4) the agent updates internal state. We run 50 sessions per model (2,500 total turns) to evaluate learning curves and convergence.

**Question pool and tier partitioning.** The main benchmark uses 20 MMLU subjects [15]; extension suites (GSM8K, HumanEval, BFCL) probe cross-task generalisation. For each model, we pre-compute per-subject accuracy on a held-out validation set and partition the 20 subjects into three




Figure 4: **CBF heatmap: 15 models  $\times$  7 baselines.** Cells colored by CBF (red–white–green = 0.40–1.00). Black-bordered cell marks the per-row argmax (= “Best” column in Table 2). Three patterns visible: (i) the algorithmic L2 cluster (UF/Platt/Bayes/EM) is uniformly high for non-frontier models but collapses to NoMem on the frontier; (ii) ZS-SK is bright green on GPT and dark red on Claude (the H2 family split); (iii) LSA dominates the Claude rows with green cells.




Figure 5: **Three reliability profiles.** Per-bin accuracy vs. confidence for three representative agents, illustrating the canonical patterns: Llama-3-8B (overconfident: bars below diagonal), Qwen-Turbo (calibrated: bars track diagonal), Claude-Opus-4-6 (conservative: bars above diagonal). Per-bin values are reconstructed from the per-model overall ECE (Table 2); the qualitative shape is empirical, the per-bin granularity is illustrative.

713 tiers *relative to the model’s own accuracy distribution*: Strong (top third,  $\geq +15$ pp above mean),  
 714 Medium (middle third,  $\pm 15$ pp), Weak (bottom third,  $\geq -15$ pp below mean). Relative partitioning  
 715 ensures every model—weak or strong—faces a non-trivial boundary: a uniform-confidence agent  
 716 cannot achieve high CBF. Per-model subject mappings are in Appendix Table 7.

717 **User personas.** The simulated user is driven by a cross-family LLM (Qwen-Max for  
 718 Qwen/Llama/Mistral agents, GPT-5.4-mini for Claude agents, Claude Sonnet 4-6 for GPT agents)  
 719 at temperature 0.7. Each persona is parameterised by an expertise set  $\mathcal{E} \subset \mathcal{D}$  (typically 5–8 of 20

720 subjects) and persona traits (direct vs. hedging style, willingness to volunteer corrections, patience).  
 721 In-expertise:  $\sim 75\%$  reliable corrections; out-of-expertise:  $\sim 46\%$  reliability with agreement-with-  
 722 wrong-answers, incorrect corrections, or vague non-commitments. This asymmetry drives the  
 723 second-order learning problem.

724 **Auxiliary metrics.** Beyond CBF (Eq. 1, computed once at the end of the 50-turn session), we  
 725 report: (i) Hallucination Rate  $HR = |\{t : c_t > 0.7 \wedge a_t \neq a_t^*\}|/T$ ; (ii) Expected Calibration Error  
 726 following Guo et al. [13] with  $M = 10$  equal-width bins,  $ECE = \sum_{m=1}^M \frac{|B_m|}{T} |\text{acc}(B_m) - \text{conf}(B_m)|$ ;  
 727 (iii) Brier Score  $= \frac{1}{T} \sum_t (c_t - \mathbb{1}[a_t = a_t^*])^2$ .

## 728 C Baseline Implementations and Design-Point Mapping

729 We instantiate ten baselines, organized by the five-level diagnostic ladder of §3.

730 **L0 — NoMemory (B1).** Raw verbalised confidence with no state. Lower bound on memory-  
 731 augmented methods.

732 **L1 — ZS-SK (B11).** The agent is prompted, *without* session history and *without* any user feedback,  
 733 to estimate its own per-domain accuracy directly: “For each of the following 20 MMLU subjects,  
 734 estimate your accuracy on a 0–100 scale.” This isolates prior metacognition.

735 **L2 — LSA (B6).** Session history (50 (Q, A, confidence, decoded-feedback) tuples) appended to  
 736 the prompt; agent reasons in-context and outputs per-domain  $\hat{p}(d)$  at the end of the session.

737 **L2 — UF-PDC (B3).** Per-domain accuracy estimate updated with labels decoded from natural-  
 738 language feedback by an LLM-based classifier (Qwen-Max, prompted to output a binary correctness  
 739 judgment for each (Q, A, feedback) triple).

740 **L2 — Platt-Online (B7) [36].** Logistic recalibration  $\sigma(wc_{\text{raw}} + b)$  with online gradient descent on  
 741 log-loss whenever decoded feedback is available; per-domain confidence is then aggregated.

742 **L2 — HistBin-Online (B8) [52].** 10-bin histogram binning with online updates; per-domain  
 743 confidence aggregated as for Platt.

744 **L2 — Bayes-Domain (B9).** Per-domain Beta( $\alpha, \beta$ ) posterior with uniform prior, updated from  
 745 decoded feedback;  $\hat{p}(d) = \alpha/(\alpha + \beta)$ .

746 **L2 — EM-Joint (B10).** Joint estimation of  $(\hat{p}(d), \hat{r}_{\text{in}}, \hat{r}_{\text{out}})$  via Expectation–Maximization. E-step  
 747 computes posterior  $P(\text{correct}_t | \text{decoded feedback}_t, p, r)$ ; M-step updates  $\hat{p}(d)$  per domain and  $\hat{r}$  per  
 748 expertise group. Initialised from raw confidence means.

749 **L3 — OraclePDC (B2).** Per-domain accuracy estimate updated with ground-truth labels after each  
 750 turn. Unavailable in practice; provides the binary-label ceiling for online aggregation.

751 **L4 — GT-SelfStats (B12).** The true per-domain accuracy  $p^*(d)$  is provided in the prompt and the  
 752 agent is asked to repeat them. Tests purely the verbalization sub-capability.

753 **Design-point mapping.** The five-level ladder is designed to subsume the design points of pub-  
 754 lished memory systems while exposing the diagnostic axis. LSA (L2) corresponds to RMM [41]  
 755 reflection over long-term memory and to in-context reasoning baselines for memory agents [16].  
 756 Platt/HistBin/Bayes (L2) cover classical online recalibration. EM-Joint (L2) tests whether explicitly  
 757 modeling the second-order coupling helps; ZS-SK (L1) and GT-SelfStats (L4) are diagnostic anchors  
 758 that isolate prior metacognition and verbalization respectively. Published systems (MemGPT [35],  
 759 Mem0 [33], A-MEM [47]) operate at L2 with different memory representations; we did not directly  
 760 integrate them because the diagnostic ladder already covers the design space. Their direct evaluation  
 761 under our protocol is engineering-feasible and is left to future work.

## D Statistical Power Analysis for the n=6 Frontier Sample

**Motivation.** The frontier-only analyses in §4.2 (DS-asym, GLAD, LSA vs. L1 ZS) use  $n = 6$  models. We document here why this sample size is adequate for the *direction-of-effect* claims we make and why it is not sufficient for absolute-magnitude claims.

**A-priori power calculation.** For a paired comparison with the observed mean improvement  $\Delta = 0.072$  (GLAD vs. L1 ZS) and the within-pair SD estimated from 10K bootstrap resamples ( $s_\Delta = 0.034$ ), the standardized effect is Cohen’s  $d_z = \Delta/s_\Delta = 1.42$ . A paired Wilcoxon signed-rank test at  $\alpha = 0.05$  (two-sided) achieves power 0.85 at  $n = 6$  for an effect of  $d_z = 1.42$  (G\*Power 3.1, exact-test option). DS-asym vs. L1 ZS has  $d_z = 1.66$  (power 0.92 at  $n = 6$ ,  $\alpha = 0.05$ ). The 6/6 sign-consistency for GLAD also independently rejects the null direction at  $\alpha = 0.016$  under a two-sided binomial sign test.

**Effect-size + bootstrap CI accompanying every n=6 claim.** For each frontier-only test we report the per-model deltas, the Cohen’s  $d_z$ , and the bootstrap 95% CI alongside any p-value:

| Comparison       | $n$ | Mean $\Delta$ | $d_z$ | Bootstrap 95% CI | Win/Loss |
|------------------|-----|---------------|-------|------------------|----------|
| GLAD vs L1 ZS    | 6   | +0.072        | 1.42  | [+0.045, +0.099] | 6/0      |
| DS-asym vs L1 ZS | 6   | +0.115        | 1.66  | [+0.058, +0.172] | 5/1      |
| LSA vs L1 ZS     | 6   | +0.108        | 0.55  | [−0.030, +0.246] | 4/2      |

LSA’s CI crosses zero (consistent with Wilcoxon  $p = 0.156$ ), so we do not claim  $\text{LSA} > \text{ZS}$  as a frontier-mean effect; the 4/2 win count reflects the family-split (Claude vs. GPT) discussed in the main text rather than a uniform effect.

**Sample selection.** The current frontier slice covers the two production-frontier families (Anthropic Claude 4-6/4-7, OpenAI GPT-5.4/5.4-mini/5.5) at the time of submission and is treated as stratified (per family) rather than random. Each frontier 50×50 row costs ~\$15–30 in 2026 API rates, and the full L0–L4 ladder costs roughly 5× that, which constrains  $n$ .

**Detection limits at  $n = 6$ .** Three comparisons that this sample does not support: (a) DS-asym vs. GLAD ( $\Delta = 0.043$  requires  $n \approx 18$  at power 0.80); (b) absolute frontier-mean estimates with 95% CI tighter than  $\pm 0.05$ ; (c) within-family contrasts at  $\alpha=0.05$ . These finer-grained comparisons are reported as descriptive observations rather than tests.

## E Additional Robustness Ablations

This appendix contains additional ablation experiments orthogonal to the main diagnostic ladder.

**A1: Memory size for L2 retrieval.** Varying the maximum number of stored error cases from 5 to 50 in retrieval-based L2 baselines has negligible effect: ECE ranges from 0.097 to 0.096 on Llama-3-8B. Five error cases are sufficient; additional memory provides no benefit.

**A2: Feedback noise rate.** Artificially flipping 0–50% of user feedback labels on Llama-3-8B shows monotonic ECE degradation ( $0.190 \rightarrow 0.238$ ) and CBF degradation. At 50% noise, learning from feedback is actively harmful (CBF 0.639 < NoMemory 0.742), confirming that noisy feedback can poison self-models.

**A3: Session length.** OraclePDC achieves near-perfect CBF ( $\geq 0.96$ ) after just 10–20 turns, while practical L2 methods plateau around 0.71 by turn 30. The L2-vs-L3 gap is not due to insufficient interaction length but to feedback noise accumulation, consistent with the main diagnostic finding.

**A4: Held-out MMLU-Pro contamination control (cited in §4.2).** For each of the 6 frontier models, we sampled 15 questions from each of 10 MMLU-Pro categories (150 questions/model, 10 categories: math, physics, chemistry, law, engineering, economics, health, psychology, business, history), measured per-category accuracy with *up to 5 retries on API failures* (persistent failures

803 excluded rather than counted as wrong), and queried ZS-SK on those category names. Per-category  
804 accuracy and ZS estimates:

| Category     | GPT-5.4      |      | GPT-5.4-mini |      | GPT-5.5      |      | Claude-Opus-4-6 |      | Claude-Sonnet-4-6 |      | Claude-Opus-4-7 |      |
|--------------|--------------|------|--------------|------|--------------|------|-----------------|------|-------------------|------|-----------------|------|
|              | acc          | ZS   | acc          | ZS   | acc          | ZS   | acc             | ZS   | acc               | ZS   | acc             | ZS   |
| math         | 0.40         | 0.85 | 0.33         | 0.78 | 0.93         | 0.72 | 0.80            | 0.78 | 0.93              | 0.72 | 0.87            | 0.78 |
| physics      | 0.40         | 0.84 | 0.47         | 0.72 | 1.00         | 0.70 | 1.00            | 0.80 | 1.00              | 0.74 | 0.73            | 0.80 |
| chemistry    | 0.60         | 0.82 | 0.20         | 0.70 | 0.87         | 0.72 | 0.93            | 0.75 | 0.92              | 0.73 | 0.73            | 0.72 |
| law          | 0.60         | 0.81 | 0.53         | 0.74 | 0.80         | 0.76 | 0.73            | 0.70 | 0.73              | 0.67 | 0.80            | 0.60 |
| engineering  | 0.53         | 0.80 | 0.47         | 0.73 | 0.80         | 0.68 | 0.93            | 0.72 | 1.00              | 0.70 | 0.67            | 0.65 |
| economics    | 0.80         | 0.83 | 0.73         | 0.79 | 0.93         | 0.78 | 1.00            | 0.84 | 0.87              | 0.76 | 0.87            | 0.82 |
| health       | 0.67         | 0.86 | 0.73         | 0.77 | 0.87         | 0.76 | 0.79            | 0.79 | 0.80              | 0.75 | 0.71            | 0.72 |
| psychology   | 0.80         | 0.88 | 0.73         | 0.82 | 0.80         | 0.82 | 0.87            | 0.87 | 0.87              | 0.78 | 0.73            | 0.85 |
| business     | 0.47         | 0.84 | 0.20         | 0.80 | 0.80         | 0.82 | 0.93            | 0.81 | 0.80              | 0.75 | 0.80            | 0.78 |
| history      | 0.73         | 0.87 | 0.67         | 0.84 | 0.80         | 0.80 | 0.80            | 0.76 | 0.73              | 0.79 | 0.73            | 0.75 |
| mean         | 0.60         | 0.84 | 0.51         | 0.77 | 0.86         | 0.77 | 0.88            | 0.78 | 0.86              | 0.74 | 0.76            | 0.75 |
| unparsed/150 | 0            |      | 0            |      | 0            |      | 1               |      | 8                 |      | 1               |      |
| CBF          | <b>0.760</b> |      | <b>0.738</b> |      | <b>0.888</b> |      | <b>0.902</b>    |      | <b>0.863</b>      |      | <b>0.941</b>    |      |

805 *Note on retry methodology:* An earlier evaluation without retries had 19–52/150 API failures counted  
806 as wrong, depressing measured accuracy and inflating CBF. We re-ran with up-to-5 retries with  
807 exponential backoff; persistent failures (0–8/150 across models) are excluded rather than counted  
808 as errors. The retry-stabilized table above is what we report. *Mean MMLU-Pro CBF across the 6*  
809 *frontier models is 0.849*, higher than the original n=4 mean of 0.816, confirming the contamination-  
810 control conclusion holds (and strengthens) on the expanded sample. The two newly added frontier  
811 models continue the pattern: GPT-5.5 (CBF 0.888) and Claude-Opus-4-7 (CBF 0.941) both report  
812 ZS estimates that are systematically conservative relative to actual MMLU-Pro accuracy (mean ZS  
813 0.77/0.75 vs. accuracy 0.86/0.76). Frontier metacognition—if anything—underestimates the model’s  
814 actual MMLU-Pro capability rather than overestimating it, the opposite of what a memorized-public-  
815 stats hypothesis would predict.

816 **A4b: Truly held-out novel-task contamination control (BBH; cited in §4.2).** For each of  
817 the 6 frontier models, we sampled 15 questions from each of 10 BBH (BIG-Bench Hard) [40]  
818 tasks (150 questions/model, tasks: *boolean\_expressions*, *causal\_judgement*, *date\_understanding*,  
819 *logical\_deduction\_five\_objects*, *navigate*, *object\_counting*, *ruin\_names*, *sports\_understanding*, *track-*  
820 *ing\_shuffled\_objects\_three\_objects*, *word\_sorting*). Unlike MMLU-Pro categories (which still par-  
821 tially overlap with MMLU’s broad subject space), the BBH task set has *no* subject-name overlap with  
822 MMLU’s 57 fine-grained subjects, providing a stricter test of whether L1 metacognition transfers  
823 beyond MMLU-style content. We measured per-task accuracy with up-to-5 retries on API failures  
824 (zero failures across all 6 frontier models, 0/900 total) and queried ZS-SK on the same task names.  
825 Per-task accuracy and ZS estimates:

| Task                 | GPT-5.4      |      | GPT-5.4-mini |      | GPT-5.5      |             | Claude-Sonnet-4-6 |      | Claude-Opus-4-6 |      | Claude-Opus-4-7 |             |
|----------------------|--------------|------|--------------|------|--------------|-------------|-------------------|------|-----------------|------|-----------------|-------------|
|                      | <i>p</i> *   | ZS   | <i>p</i> *   | ZS   | <i>p</i> *   | ZS          | <i>p</i> *        | ZS   | <i>p</i> *      | ZS   | <i>p</i> *      | ZS          |
| boolean_expressions  | 0.87         | 0.94 | 0.87         | 0.82 | 1.00         | 0.95        | 1.00              | 0.88 | 0.93            | 0.96 | 0.87            | 0.95        |
| causal_judgement     | 0.67         | 0.83 | 0.67         | 0.78 | 0.67         | 0.85        | 0.73              | 0.60 | 0.53            | 0.72 | 0.80            | 0.65        |
| date_understanding   | 0.27         | 0.90 | 0.20         | 0.75 | 0.47         | 0.85        | 1.00              | 0.82 | 0.93            | 0.92 | 0.87            | 0.85        |
| logical_deduction_5  | 1.00         | 0.86 | 0.67         | 0.70 | 1.00         | 0.90        | 1.00              | 0.65 | 1.00            | 0.88 | 1.00            | 0.90        |
| navigate             | 0.80         | 0.88 | 0.53         | 0.68 | 1.00         | 0.85        | 1.00              | 0.76 | 0.93            | 0.95 | 0.93            | 0.90        |
| object_counting      | 0.73         | 0.96 | 0.53         | 0.90 | 1.00         | 0.90        | 1.00              | 0.84 | 0.80            | 0.97 | 0.60            | 0.95        |
| ruin_names           | 0.67         | 0.79 | 0.67         | 0.72 | 0.93         | 0.75        | 0.87              | 0.56 | 0.87            | 0.86 | 0.93            | 0.75        |
| sports_understanding | 1.00         | 0.84 | 0.93         | 0.74 | 0.80         | 0.90        | 1.00              | 0.90 | 0.93            | 0.96 | 0.87            | 0.90        |
| tracking_shuffled_3  | 0.67         | 0.95 | 0.67         | 0.88 | 1.00         | 0.85        | 1.00              | 0.73 | 1.00            | 0.98 | 0.53            | 0.95        |
| word_sorting         | 0.67         | 0.97 | 0.47         | 0.95 | <b>1.00</b>  | <b>0.90</b> | 0.00              | 0.80 | 0.13            | 0.94 | <b>0.93</b>     | <b>0.90</b> |
| CBF                  | <b>0.781</b> |      | <b>0.780</b> |      | <b>0.850</b> |             | <b>0.734</b>      |      | <b>0.861</b>    |      | <b>0.860</b>    |             |

826 *Mean BBH CBF across all 6 frontier models is 0.811; under the 50×50 protocol mean MMLU ZS*  
827 *CBF for the same 6 models is 0.640*, so BBH is +0.171 higher than MMLU—the opposite of an  
828 MMLU-memorization prediction, and the gap widens slightly when extended from n=4 (+0.165)  
829 to n=6. The two newly added frontier models (GPT-5.5 BBH 0.850, Claude-Opus-4-7 BBH 0.860)  
830 both exceed the n=4 mean (0.789) and place near the top of the BBH ranking, while their MMLU ZS

values (GPT-5.5 0.859, Opus-4-7 0.482) span the full range. *Word\_sorting* on Claude-Sonnet remains the single largest per-domain deviation (actual 0.00, ZS 0.80). Four findings stand out:

(1) *word\_sorting* blind-spot is generation-specific, not LLM-wide. The original 4 frontier models all substantially overestimate *word\_sorting* (actuals 0.00–0.67, ZS estimates 0.80–0.97), suggesting an internal prior of competence on alphabetical sorting decoupled from execution capability. The 2 newly added frontier models break this pattern: GPT-5.5 reaches 1.00 (ZS 0.90, slight underestimate) and Claude-Opus-4-7 reaches 0.93 (ZS 0.90, well-calibrated). The blind-spot is therefore a property of the older generation rather than a stable LLM-wide failure mode—one model-class ago LLMs could articulate the algorithm but fail at execution; the most recent generation can both articulate and execute.

(2) *date\_understanding* remains a GPT-side overestimate. GPT-5.4, GPT-5.4-mini, and GPT-5.5 all substantially overestimate this task (actuals 0.20–0.47 vs. ZS 0.75–0.90). The Claude models are well-calibrated here (Sonnet 1.00/0.82, Opus-4-6 0.93/0.92, Opus-4-7 0.87/0.85). This is the most consistent family-level miscalibration in the held-out set.

(3) Opus-4-7 has new failure modes on object/tracking tasks. Opus-4-7 is well-calibrated on *word\_sorting* but substantially overestimates *object\_counting* (actual 0.60 vs. ZS 0.95) and *tracking\_shuffled\_3* (actual 0.53 vs. ZS 0.95). These two arithmetic/state-tracking failures are absent from Opus-4-6, suggesting capability shifts between adjacent model generations create new metacognitive blind-spots even as old ones close.

(4) Outside the family-specific outliers, BBH ZS-SK is well-calibrated across all 6 models. If we exclude *word\_sorting* and *date\_understanding*, mean ZS error on the remaining 8 tasks is 0.13 (GPT-5.4), 0.16 (GPT-5.4-mini), 0.13 (GPT-5.5), 0.20 (Claude-Sonnet-4-6), 0.07 (Claude-Opus-4-6), 0.17 (Claude-Opus-4-7). The five non-outlier model means correspond to per-domain CBFs of 0.83–0.93, comparable to or exceeding MMLU performance.

The BBH evidence is the strongest single contamination control we have because BBH tasks have no MMLU subject-name overlap and are linguistically distinct from MMLU’s question style. The result—mean BBH CBF 0.811 across all 6 frontier models—directly contradicts what an MMLU-memorization explanation would predict (a noticeable drop on novel content). Combined with the MMLU-Pro result, we conclude that L1 metacognition reflects genuine prior self-knowledge, with identifiable family- and generation-specific failure modes (*word\_sorting* for the older 4, *date\_understanding* for GPT, *object\_counting* and *tracking\_shuffled* for Opus-4-7) that point to a follow-up direction: building an internal task-difficulty model that is independent of the model’s self-perceived competence on similar-sounding tasks.

**A5: Stronger crowd-sourcing estimators (cited in §4.3).** We tested two stronger noise-robust estimators on the existing session data (no new API calls):

**Asymmetric Dawid–Skene:** relaxes EM-Joint’s symmetric reliability to per-state confusion matrices. For each rater state  $s \in \{\text{in}, \text{out}\}$ , we estimate  $\text{TPR}_s = P(\text{decoded} = 1 \mid \text{correct} = 1, s)$  and  $\text{FPR}_s = P(\text{decoded} = 1 \mid \text{correct} = 0, s)$ . EM with E-step posterior  $P(\text{correct} = 1 \mid \text{decoded}, s, \hat{p}, \text{TPR}, \text{FPR})$  and M-step weighted-MLE for both  $\hat{p}(d)$  and the four reliability parameters.

**GLAD-style:** per-domain difficulty  $\delta_d$  and per-state ability  $\alpha_s$ , with reliability  $\sigma(\alpha_s - \delta_d)$  (logistic). EM updates over  $\hat{p}(d)$ ,  $\delta_d$ ,  $\alpha_{\text{in}}$ ,  $\alpha_{\text{out}}$ .

Per-model CBF results:

| Model                                  | EM-Joint | DS-asym      | GLAD         | ZS-SK (L1) |
|----------------------------------------|----------|--------------|--------------|------------|
| Mistral-7B                             | 0.722    | 0.706        | 0.693        | 0.678      |
| Llama-3-8B                             | 0.716    | 0.646        | 0.612        | 0.682      |
| GLM-5                                  | 0.768    | 0.615        | 0.683        | 0.590      |
| DeepSeek-V3.2                          | 0.708    | 0.702        | 0.738        | 0.699      |
| Qwen-Turbo                             | 0.731    | 0.714        | 0.776        | 0.751      |
| Qwen2.5-7B                             | 0.760    | 0.774        | 0.764        | 0.684      |
| Qwen-Max                               | 0.702    | 0.547        | 0.627        | 0.767      |
| Qwen-Plus                              | 0.712    | 0.762        | 0.761        | 0.740      |
| Qwen3.5-27B                            | 0.751    | 0.719        | 0.733        | 0.695      |
| GPT-5.4-mini                           | 0.706    | 0.797        | <b>0.797</b> | 0.769      |
| GPT-5.4                                | 0.677    | 0.815        | <b>0.843</b> | 0.795      |
| GPT-5.5                                | 0.605    | 0.834        | <b>0.909</b> | 0.859      |
| Claude-Sonnet-4-6                      | 0.609    | <b>0.671</b> | 0.565        | 0.472      |
| Claude-Opus-4-6                        | 0.604    | <b>0.734</b> | 0.557        | 0.460      |
| Claude-Opus-4-7                        | 0.570    | <b>0.676</b> | 0.600        | 0.482      |
| <b>Mean (n=15)</b>                     | 0.689    | <b>0.714</b> | 0.711        | 0.675      |
| <b>Frontier mean (n=6)<sup>‡</sup></b> | 0.628    | <b>0.755</b> | 0.712        | 0.640      |

Under the 50×50 protocol, the strongest second-order estimator on frontier is Dawid-Skene-asymmetric (mean CBF 0.755, n=6, Wilcoxon one-sided  $p = 0.047$  vs L1); GLAD reaches mean CBF 0.712 (Wilcoxon two-sided  $p = 0.031$ , 6/6 wins, binomial one-sided  $p = 0.016$ ), beating L1 ZS mean 0.640 by +0.072 for GLAD and +0.115 for DS. Per-frontier breakdown (GLAD vs L1 ZS): GPT-5.4 +0.048, GPT-5.4-mini +0.028, GPT-5.5 +0.050, Claude-Sonnet-4-6 +0.093, Claude-Opus-4-6 +0.097, Claude-Opus-4-7 +0.118. GLAD now beats L1 on all 6/6 frontier models (under 50×50, GPT-5.5 GLAD is 0.909 vs ZS 0.859). DS-asym beats L1 on 5/6 (loses only on GPT-5.5 by −0.025, where GLAD’s per-domain difficulty modeling helps). *Methodological note:* our DS / GLAD implementations follow Dawid-Skene (1979) and Whitehill et al. (2009) respectively, with EM updates over per-domain  $\hat{p}(d)$ , asymmetric per-state confusion (DS) or per-domain difficulty  $\delta_d$  + per-state ability  $\alpha_s$  (GLAD); convergence within 50–80 EM iterations using initial  $\hat{p} = 0.5$ ,  $r_{\text{in}} = 0.75$ ,  $r_{\text{out}} = 0.55$ . Code at the anonymized repository/scripts/run\_glad\_ds.py. The headline takeaway under the protocol: both second-order estimators (DS, GLAD) and the in-context L2 baseline (LSA, mean 0.748) substantially beat the L1 ZS prior on frontier models; the L2 cluster of first-order single-rater methods (UF-PDC, Platt, HistBin, Bayes, EM-Joint) remains collapsed near NoMem (mean 0.628–0.634). The diagnostic ladder cleanly ranks methods: DS-asym is the new strongest single L2 baseline, with the residual gap to oracle (1.000) the open problem.

**A6: Decoder ablation for EM-Joint (cited in §4.3).** We isolate the EM machinery from the feedback decoder by running EM-Joint with three different label channels:

| Model                          | Oracle-fed | LLM-decoder | Rule-decoder |
|--------------------------------|------------|-------------|--------------|
| Mistral-7B                     | 1.000      | 0.722       | 0.672        |
| Llama-3-8B                     | 1.000      | 0.716       | 0.549        |
| GLM-5                          | 1.000      | 0.768       | 0.672        |
| DeepSeek-V3.2                  | 1.000      | 0.708       | 0.690        |
| Qwen-Turbo                     | 1.000      | 0.731       | 0.749        |
| Qwen2.5-7B                     | 1.000      | 0.760       | 0.676        |
| Qwen-Max                       | 1.000      | 0.702       | 0.771        |
| GPT-5.4-mini                   | 1.000      | 0.632       | 0.633        |
| Qwen-Plus                      | 1.000      | 0.712       | 0.758        |
| Qwen3.5-27B                    | 1.000      | 0.751       | 0.756        |
| GPT-5.4                        | 1.000      | 0.613       | 0.673        |
| Claude-Sonnet-4-6              | 1.000      | 0.608       | 0.896        |
| Claude-Opus-4-6                | 1.000      | 0.644       | 0.885        |
| <b>Mean (n=13)<sup>b</sup></b> | 1.000      | 0.698       | 0.715        |

<sup>b</sup>This decoder ablation was conducted on the 13 models available before the GPT-5.5 and Claude-Opus-4-7 frontier additions; extending the EM oracle/decoder ablation to the two newly-added rows




Figure 6: **Learning dynamics:** CBF over interaction turns for several L2 baselines on a weak model (Llama-3-8B, left) and a strong model (GPT-5.4, right). On both models the L2 methods plateau well below the OraclePDC ceiling, with no method monotonically improving past  $\sim 0.72$  on either model. NoMemory CBF is approximately constant as expected (no learning occurs). This figure is illustrative; the headline diagnostic in the main paper uses end-of-session CBF (Eq. 1).

897 is left to future work and does not change the qualitative finding (the LLM-decoder vs. rule-decoder  
898 gap to oracle, see below).

899 Three observations: (a) *Oracle-fed EM* reaches  $\text{CBF} = 1.000$  on every model, identical to  
900 OraclePDC—this confirms the EM update equations themselves are not the bottleneck; if per-  
901 fect labels are available, EM reduces to per-domain counting and trivially recovers  $p^*(d)$ . (b) Two  
902 structurally very different feedback decoders—an LLM-based binary classifier (Qwen-Max prompted  
903 with the (Q, A, feedback) triple) and a regex-rule-based decoder (matching keywords like “correct”,  
904 “actually”, “wrong”)—both produce noisy estimates that fall *far below* oracle (means 0.698 vs. 0.715).  
905 (c) The two decoders disagree per-model in opposite directions for some agents (Claude-Opus: rule  
906  $0.885 > \text{LLM } 0.644$ ; Llama-3-8B:  $\text{LLM } 0.716 > \text{rule } 0.549$ ), indicating they capture different subsets  
907 of the noise pattern. The persistent gap from oracle (1.000) to either decoder (mean  $\approx 0.71$ ) is  
908 therefore a property of the *noisy feedback channel itself*, not of any particular decoding strategy.  
909 Stronger decoders or crowd-sourcing-style estimators (Dawid-Skene, GLAD) might close some  
910 of this gap on a per-model basis, but in our setting the residual is primarily information-theoretic:  
911 per-turn evidence about  $p^*(d)$  from natural-language user feedback is sparse and domain-correlated.

## 912 F Learning Dynamics

913 Figure 6 shows the learning dynamics. CBF improves most rapidly during the first 100–150 turns  
914 before plateauing. For NoMemory, CBF remains approximately constant. For L2 feedback-driven  
915 methods, CBF typically reaches its best value around turn 200 and then plateaus or slightly degrades  
916 due to accumulated noise from decoded feedback labels. For GPT-5.4 (strong model), the noise in  
917 decoded feedback consistently fails to improve a model whose prior is already well-differentiated  
918 across domains—consistent with the main paper’s diagnostic finding that L2 methods do not exceed  
919 the L1 prior on frontier models. This suggests that practical methods may benefit from a forget-  
920 ting mechanism—discounting older feedback as the agent’s self-model stabilizes—analogueous to  
921 exponential forgetting in adaptive filtering [21].

## 922 G Statistical Robustness

923 We assess stability via bootstrap resampling using 200 sessions resamples. The table below reports  
924 NoMemory baseline estimates with bootstrap means  $\pm$  SDs across 13 models:

| Model             | NoMem CBF | ECE       | HR        | Brier     |
|-------------------|-----------|-----------|-----------|-----------|
| Mistral-7B        | .691±.010 | .585±.020 | .585±.025 | .582±.022 |
| Llama-3-8B        | .634±.008 | .446±.025 | .514±.027 | .449±.022 |
| GLM-5             | .701±.008 | .228±.009 | .018±.005 | .123±.006 |
| DeepSeek-V3.2     | .647±.011 | .158±.014 | .092±.010 | .139±.009 |
| Qwen-Turbo        | .654±.009 | .198±.013 | .203±.017 | .201±.012 |
| Qwen2.5-7B        | .695±.013 | .256±.020 | .350±.021 | .284±.014 |
| Qwen-Max          | .654±.014 | .182±.015 | .198±.013 | .190±.013 |
| GPT-5.4-mini      | .576±.005 | .092±.011 | .121±.011 | .105±.009 |
| Qwen-Plus         | .648±.010 | .177±.010 | .163±.009 | .166±.008 |
| Qwen3.5-27B       | .679±.014 | .174±.009 | .026±.007 | .102±.007 |
| GPT-5.4           | .561±.008 | .069±.009 | .086±.009 | .079±.007 |
| Claude-Sonnet-4-6 | .553±.005 | .085±.009 | .026±.005 | .059±.007 |
| Claude-Opus-4-6   | .572±.007 | .066±.008 | .034±.007 | .052±.005 |

The bootstrap SDs are consistently small (CBF  $\pm 0.005$ – $0.014$ , ECE  $\pm 0.008$ – $0.025$ ), confirming our findings are stable under session-level resampling. The NoMemory bootstrap means are within 1–3 SDs of the corresponding values in Table 2 ; the small offset reflects bootstrap shrinkage and is consistent across models. Bootstrap SDs for the other baselines (LSA, EM-Joint, etc.) are of comparable magnitude (typically  $\leq 0.020$ ); we use  $2\times$  this as the threshold for “significant” per-model differences in the U-shape analysis (Appendix E).

**Hierarchical bootstrap for cross-model means.** The session-level resampling above quantifies within-model uncertainty but does not capture variation across the 15-model sample. We additionally implement a hierarchical bootstrap that resamples models (with replacement) and, within each sampled model, resamples sessions (with replacement); CBF is recomputed on each resample. We run 1,000 hierarchical iterations to derive 95% CIs for cross-model mean CBF. Two findings are worth highlighting:

Table 4: Pairwise CBF mean differences across all 15 models under the  $50\times 50$  protocol (point estimates from Table 2; the hierarchical bootstrap CIs are based on the protocol). All cross-model differences in this table are within  $\pm 0.04$ , confirming that no single L0–L2 baseline strictly dominates at the population level.

| Comparison              | Mean diff (raw CBF, n=15) |
|-------------------------|---------------------------|
| ZS-SK – NoMemory        | +0.014                    |
| ZS-SK – LSA             | −0.017                    |
| LSA – NoMemory          | +0.032                    |
| EM-Joint – NoMemory     | +0.015                    |
| Platt-Online – NoMemory | +0.006                    |

Table 5: Frontier (n=6) vs. non-frontier (n=9) mean CBF difference by baseline under the  $50\times 50$  protocol. All baselines except LSA are lower on the frontier; LSA is the only baseline whose mean rises on the frontier subset, which is the key non-monotonicity that motivates the H2/H3 family-split discussion in §4.2.

| Baseline       | Frontier – Non-frontier |
|----------------|-------------------------|
| NoMemory       | −0.101                  |
| ZS-SK          | −0.076                  |
| LSA            | +0.076                  |
| UF-PDC         | −0.065                  |
| Platt-Online   | −0.079                  |
| HistBin-Online | −0.079                  |
| Bayes-Domain   | −0.093                  |
| EM-Joint       | −0.095                  |

Interpretation (updated for  $50\times 50$  protocol). Three observations from the point-estimate tables. *First*, the population-level mean improvements on n=15 are all small ( $\leq +0.032$ ): no L0–L2 baseline strictly dominates the others by a large cross-model margin, and all pairwise differences are within

$\pm 0.04$ . *Second*, LSA is the only baseline whose mean *rises* on the frontier subset ( $n=6$  LSA mean 0.748 vs.  $n=9$  mean 0.672 =  $+0.076$ ); every other baseline (NoMem, ZS, all algorithmic L2) shows lower mean CBF on the frontier than non-frontier. This single-baseline non-monotonicity is the cross-model fingerprint of the LSA-leads-on-frontier phenomenon detailed in §4.2. *Third*, ZS–NoMem (population mean  $+0.014$ ) is small overall but the frontier breakdown is family-specific: ZS exceeds NoMem on all 3 GPT frontier models but falls below NoMem on all 3 Claude frontier models, so the population mean averages a strong-positive GPT effect with a strong-negative Claude effect. This is consistent with the H2 family-split finding. The full hierarchical bootstrap CIs for all of the above are reported as point estimates here for the protocol.

**Hierarchical-vs-session-only CI width.** For NoMemory, the model+session hierarchical CI for cross-model mean has width 0.107, vs. 0.072 for the model-level-only bootstrap—confirming that ignoring within-model session variance underestimates uncertainty by  $\approx 1.5\times$ . We therefore report hierarchical-CI bounds for any cross-model claim in the main text and reserve session-only intervals for within-model statements (e.g. U-shape per-model significance in Appendix E).

## H Normalized Hallucination Rate

Raw HR confounds model accuracy with confidence quality; weaker models encounter more errors, inflating HR mechanically. We report *normalized* HR = HR / base error rate, where 1.0 indicates no better than random high-confidence assignment:

| Model      | Base Err | Raw HR | Norm HR |
|------------|----------|--------|---------|
| Mistral-7B | 0.596    | 0.590  | 0.990   |
| Llama-3-8B | 0.448    | 0.448  | 1.000   |
| Qwen-Turbo | 0.268    | 0.221  | 0.824   |
| Qwen-Max   | 0.244    | 0.206  | 0.846   |

Llama-3-8B and Mistral-7B have normalized HR  $\approx 1.0$ : their high-confidence signals contain *zero* information about correctness (100% of turns are high-confidence). Qwen models achieve normalized HR  $< 0.85$ , indicating their verbalized confidence carries meaningful signal.

## I CBF vs. Rank Correlation

A natural alternative to CBF is Kendall’s  $\tau$  over domain accuracy rankings. However,  $\tau$  is undefined when the baseline produces constant predictions; e.g., NoMemory estimates 0.5 for all domains. CBF avoids this by measuring absolute deviation rather than rank correlation, making it applicable to all baselines including the constant-prediction case. We recommend reporting both CBF and  $\tau$  when applicable.

## J Learning-Curve Saturation

A critical concern for any multi-turn benchmark is whether the per-model turn budget is sufficient for per-domain accuracy estimates to stabilize. The learning-curve analysis below was conducted on the  $10\times 50 = 500$ -turn-per-model checkpoints; the  $50\times 50$  protocol uses 2,500 turns per model and the qualitative finding (estimates stabilize by turn 200–300, with  $\leq 4\%$  deviation by turn 400) is robust to that change. We compute per-domain accuracy at checkpoints  $t \in \{50, 100, 200, 300, 400, 500\}$  and measure the mean absolute deviation from the final turn-500 estimate:

| Model           | T=50  | T=100 | T=200 | T=300 | T=400 |
|-----------------|-------|-------|-------|-------|-------|
| Llama-3-8B      | 0.160 | 0.134 | 0.081 | 0.069 | 0.035 |
| Qwen-Plus       | 0.188 | 0.099 | 0.055 | 0.060 | 0.025 |
| Qwen-Max        | 0.147 | 0.144 | 0.093 | 0.070 | 0.035 |
| Claude-Opus-4-6 | 0.096 | 0.091 | 0.065 | 0.042 | 0.027 |
| GPT-5.4         | 0.068 | 0.063 | 0.042 | 0.038 | 0.022 |

At turn 400 (within the 500-turn checkpoint range), the mean deviation from final is 2–4% across all five models, indicating reasonable convergence; at turn 200 the deviation remains 4–9%, so even the




Figure 7: Three-way relationship: MMLU accuracy ( $x$ ), best ECE ( $y$ ), and NoMemory CBF (bubble size). **Weak models** (Mistral, Llama-3) have large bubbles (high  $\text{CBF} \approx 0.84$ —close to the trivial  $\hat{p} = 0.5$  baseline) but scattered ECE. **Frontier models** (Claude, GPT) have small bubbles (low  $\text{CBF} \approx 0.58$ —their accuracy varies widely across domains, far from 0.5) despite decent ECE. This visualizes why CBF and ECE are non-redundant: CBF penalizes models with heterogeneous domain accuracy even if they are well-calibrated per-query.

conservative 500-turn budget provided meaningful stability. Under the  $50 \times 50$  protocol (2,500 turns per model) the convergence margin is correspondingly looser, and the NoMemory CBF estimate varies by only  $\pm 0.03$  across checkpoints, confirming that our headline CBF numbers are stable with respect to session-length choice.

## K Noise Sensitivity

Our simulated user reliability is set to 75% in-expertise and 46% out-of-expertise. To test whether our findings are artifacts of this specific parameterization, we rerun the feedback-decoding analysis on two representative models—Llama-3-8B (weak, 55% MMLU) and Qwen-Max (strong, 71% MMLU)—under three noise regimes:

| Model          | Regime     | In/Out | UF-PDC ECE | UF CBF |
|----------------|------------|--------|------------|--------|
| Llama-3 (55%)  | Low noise  | 90/70  | 0.044      | 0.860  |
|                | Current    | 75/46  | 0.134      | 0.783  |
|                | High noise | 60/30  | 0.184      | 0.718  |
| Qwen-Max (71%) | Low noise  | 90/70  | 0.134      | 0.848  |
|                | Current    | 75/46  | 0.218      | 0.795  |
|                | High noise | 60/30  | 0.310      | 0.696  |

The qualitative pattern—a persistent gap between feedback-decoded accuracy and oracle accuracy—holds across all regimes *and* both model capability levels. Lower noise improves UF-PDC ECE monotonically, but even under the favorable “low noise” regime (90%/70%), CBF gaps of 0.14 (Llama-3) and 0.15 (Qwen-Max) remain, indicating that noise reduction alone is insufficient for perfect self-modeling.

## L Cross-Baseline Bootstrap Significance

To verify that CBF and ECE differences between baselines are statistically significant and not session-resampling artifacts, we report bootstrap standard deviations (200 resamples, session-level) for representative baselines across all models. Typical CBF SDs are 0.012–0.021 and ECE SDs are 0.008–0.018, meaning that differences of  $\geq 0.04$  (CBF) or  $\geq 0.03$  (ECE) are significant at the 2-SD level. The headline diagnostic gaps under the  $50 \times 50$  protocol—LSA over ZS-SK by 0.248 on Claude-Sonnet-4-6 (LSA 0.720 vs. ZS 0.472), OraclePDC over EM-Joint by 0.31 on average across  $n=15$  (Oracle 1.000 vs. EM 0.686)—are all  $> 5 \times$  the bootstrap SD. With 15 models  $\times$  10 baselines = 130 pairwise comparisons, a Bonferroni correction would raise the significance threshold from  $2\sigma$  to  $\sim 3.5\sigma$  ( $\Delta \geq 0.06$ ). Under this stricter criterion, the three primary diagnostic findings (H1 ruled out as the limiting factor, H2 splits cleanly along model family with GPT-frontier ZS competitive but Claude-frontier ZS inflated, H3 localized: first-order single-rater methods cluster near NoMem while second-order DS-asym/GLAD and in-context LSA exceed L1 ZS on the frontier) all remain comfortably significant.

## M Dual-Role Bias

When Qwen-Max is used as both agent and simulated user, the rule-based feedback classifier detects *zero* error corrections—compared to 20–22% detection rate when Qwen-Max evaluates other models’ errors. Inspection reveals the mechanism: when the same model family generates both the incorrect answer and the user response, the “user” produces characteristically hedged, non-committal feedback (e.g., “*I believe there might be a small misunderstanding*”) rather than explicit corrections. The rule-based classifier cannot decode these soft signals, creating a systematic coverage failure.

Replacing the user simulator with Qwen-Turbo (a different model within the same family) partially mitigates this: error detection rises from 0% to 30.4%. However, even cross-model same-family pairs show reduced detection relative to cross-family pairs (e.g., Qwen-Max evaluating Llama-3 achieves 20.5%). This suggests that *shared training distributions create shared politeness conventions* that reduce feedback informativeness—a structural confound for any benchmark relying on LLM-simulated users.

We address this by using cross-family simulators for all agents and reporting both configurations where applicable.

## N Cross-Simulator Ablation

To verify that benchmark results are not artifacts of the specific simulator choice, we compare Qwen-Max agent evaluated with two different user simulators: Qwen-Max (same-family) and Qwen-Turbo (cross-model, same family):

| Simulator                 | Acc   | ECE   | HR    | CBF   |
|---------------------------|-------|-------|-------|-------|
| user = Qwen-Max (same)    | 79.2% | 0.137 | 20.6% | 0.708 |
| user = Qwen-Turbo (cross) | 78.9% | 0.142 | 20.9% | 0.717 |
| $ \Delta $                | 0.3pp | 0.005 | 0.3pp | 0.009 |

The differences are negligible:  $\Delta ECE = 0.005$  (well within the bootstrap SD of 0.009 for Qwen-Max),  $\Delta CBF = 0.009$ . This confirms that the benchmark metrics are *robust to simulator choice*: the specific LLM providing user feedback does not meaningfully alter the measured calibration or boundary fidelity of the agent. The dual-role bias documented in Section 4.3 affects feedback classifier coverage (0% vs. 30.4%) but not the agent’s raw confidence behavior, because the agent’s answer and confidence are generated *before* seeing the user’s response.

## O R4 Adaptive Selection and Second-Order Estimators: Full Details

R4 is the parameter-free per-model selection rule introduced in §4.3: select L1 (ZS) when  $ZS(m) > \text{NoMem}(m)$ , otherwise select L2 (LSA). This appendix provides the per-model breakdown, sign-test arithmetic, LOOCV details, and a per-frontier-model visualization of the second-order estimators.




Figure 8: Second-order estimators break the L1 prior on frontier models. Per-frontier-model CBF for L0 NoMem, L1 ZS-SK, L2 GLAD, and L2 DS-asym. **GLAD beats L1 ZS on all 6/6 frontier models** (Wilcoxon two-sided  $p = 0.031$ ; binomial one-sided  $p = 0.016$ ; Cohen’s  $d_z = 1.42$ ); DS-asym beats L1 on 5/6 (loses only on GPT-5.5 to GLAD by  $-0.025$ ). The  $n=6$  frontier mean delta is  $+0.072$  for GLAD and  $+0.115$  for DS-asym; bootstrap 95% CI of the mean is  $[+0.045, +0.099]$  for GLAD (Appendix D). On Claude rows where ZS is inflated relative to true accuracy, second-order estimators recover the signal far better than ZS or first-order L2.

1038 **Per-model R4 selections and outcomes.** R4 chooses ZS on 6/15 models (where ZS strictly exceeds  
1039 NoMem: Qwen-Turbo, Qwen-Max, Qwen-Plus, GPT-5.4-mini, GPT-5.4, GPT-5.5) and LSA on  
1040 9/15 (the weak-prior or Claude cases where NoMem ties or exceeds ZS: Mistral-7B, Llama-3-8B,  
1041 GLM-5, DeepSeek-V3.2, Qwen2.5-7B, Qwen3.5-27B, plus Claude-Sonnet-4-6, Claude-Opus-4-6,  
1042 Claude-Opus-4-7). R4 correctly avoids LSA’s catastrophic failure on Qwen-Plus (LSA 0.435 vs ZS  
1043 0.735) by selecting ZS, but misses the GPT-5.5 case where LSA 0.902 actually exceeds ZS 0.859 (R4  
1044 selects ZS, costing  $-0.043$ ).

1045 **Sign-test arithmetic.** R4 attains mean CBF 0.722 across all 15 models, exceeding always-ZS by  
1046  $+0.037$  in mean and always-LSA by  $+0.020$  in mean. Per-model sign tests, however, do *not* reach  
1047 conventional significance: **R4 vs. always-ZS: 5 wins, 1 loss, 9 ties** (effective  $n = 6$  after excluding  
1048 R4-picks-ZS ties; two-sided binomial  $p = 0.22$ ); **R4 vs. always-LSA: 4 wins, 2 losses, 9 ties** ( $n = 6$ ;  
1049  $p = 0.69$ ). The mean improvement is positive while the sign tests are non-significant because R4’s  
1050 gains are concentrated in a few models with large magnitudes (the three Claude models on R4-vs-ZS,  
1051 where LSA outperforms ZS by  $+0.179$  to  $+0.248$  and R4 correctly switches to LSA), while its  
1052 losses are smaller in count but include a single large hit (Llama-3-8B on R4-vs-ZS, where LSA  
1053 underperforms ZS by  $-0.162$  and R4 still picks LSA because ZS = NoMem on that row). We report  
1054 the sign test rather than a Wilcoxon signed-rank because the deltas are bimodal (large Claude wins vs.  
1055 a single large Llama-3 loss), violating the Wilcoxon rank assumption of unimodal symmetric deltas.

1056 **LOOCV.** When learning the threshold  $\theta$  from the remaining 14 models for each held-out model, R4  
1057 LOOCV mean is 0.719 (vs in-sample 0.722, overfit gap only  $+0.003$ ), confirming the parameter-free  
1058  $\theta = 0$  rule generalizes. The per-model oracle (best of {ZS, LSA}) attains 0.736; R4 closes 71.5% of  
1059 the always-ZS-to-oracle gap.

1060 **Closure-rate arithmetic.** R4 mean  $0.7217$  minus ZS mean  $0.685 = +0.037$ , divided by oracle-ZS  
1061 gap  $0.736 - 0.685 = 0.051$ , gives  $0.037/0.051 \approx 0.725$  (rounded to “ $\approx 72\%$ ” throughout the paper).

| Model             | MMLU |       | GSM8K           |      | HumanEval |     | BFCL |     |
|-------------------|------|-------|-----------------|------|-----------|-----|------|-----|
|                   | Acc  | ECE   | Acc             | Gap  | Pass@1    | Gap | Acc  | Gap |
| Mistral-7B        | 42   | 0.036 | 55              | +31  | 31        | +52 | 71   | +15 |
| Llama-3-8B        | 55   | 0.167 | 84              | −3   | 33        | +44 | 46   | +36 |
| GLM-5             | 55   | 0.213 | 96              | +1   | 91        | −49 | 64   | +32 |
| DeepSeek-V3.2     | 62   | 0.158 | 94              | −21  | 87        | −30 | 75   | +5  |
| Qwen-Turbo        | 64   | 0.165 | 95              | −33  | 74        | −1  | 72   | +22 |
| Qwen2.5-7B        | 65   | 0.254 | 90              | −18  | 64        | −4  | 73   | +18 |
| Qwen-Max          | 71   | 0.155 | 97              | −18  | 82        | −31 | 74   | +17 |
| Qwen-Plus         | 79   | 0.155 | 90              | −11  | 92        | −48 | 80   | +10 |
| Qwen3.5-27B       | 79   | 0.156 | 98              | −1   | 77        | −35 | 72   | +24 |
| GPT-5.4-mini      | 72   | 0.211 | 95              | −2   | 74        | −25 | 76   | +15 |
| GPT-5.4           | 82   | 0.133 | 97              | +0.3 | 81        | −36 | 70   | +25 |
| Claude-Sonnet-4-6 | 86   | 0.211 | 88              | −4   | 79        | −40 | 72   | +17 |
| Claude-Opus-4-6   | 90   | 0.158 | 99              | −5   | 74        | −17 | 78   | +9  |
| GPT-5.5           | 94   | 0.036 | 97              | −1   | 85        | −31 | 97   | −4  |
| Claude-Opus-4-7   | 69   | 0.170 | 46 <sup>b</sup> | −2   | 95        | −34 | 99   | −9  |

Table 6: Unified comparison across four task types (all 15 models, units in %). MMLU columns report NoMemory-baseline Acc and ECE; the other three task columns report accuracy and Gap (= mean confidence − accuracy; positive = overconfident, negative = conservative). Two key observations: (i) the Gap sign for the same model can flip across tasks (e.g. Llama-3-8B is conservative on GSM8K −3pp but severely overconfident on HumanEval/BFCL +44/+36pp); (ii) strong models are not uniformly conservative (Claude-Opus is conservative on GSM8K/HumanEval −5/−17pp but still overconfident on BFCL +9pp).

## P Cross-Task Heterogeneity: Full Per-Task Comparison

## Q Lite Protocol Feasibility and Controlled-Protocol Cross-Task Subset

### Q.1 Controlled-protocol cross-task subset (referenced in §5)

The cross-task heterogeneity finding in §5 compares MMLU (50-turn noisy-feedback protocol) with GSM8K, HumanEval, and BFCL (single-shot two-step protocol). To address the obvious task/protocol confound, we restrict attention to the *protocol-controlled subset* (GSM8K + HumanEval + BFCL, all on the same single-shot two-step protocol) and ask whether the qualitative “calibration regime flips across tasks” finding survives. Counting per-model sign flips of the confidence–accuracy Gap across the three same-protocol tasks (data from Table 6, columns GSM8K-Gap, HumanEval-Gap, BFCL-Gap):

| Model                                             | GSM8K Gap | HumanEval Gap | BFCL Gap | Sign-flip in subset? |
|---------------------------------------------------|-----------|---------------|----------|----------------------|
| Mistral-7B                                        | +31       | +52           | +15      | no (all +)           |
| Llama-3-8B                                        | −3        | +44           | +36      | <b>yes</b>           |
| GLM-5                                             | +1        | −49           | +32      | <b>yes</b>           |
| DeepSeek-V3.2                                     | −21       | −30           | +5       | <b>yes</b>           |
| Qwen-Turbo                                        | −33       | −1            | +22      | <b>yes</b>           |
| Qwen2.5-7B                                        | −18       | −4            | +18      | <b>yes</b>           |
| Qwen-Max                                          | −18       | −31           | +17      | <b>yes</b>           |
| Qwen-Plus                                         | −11       | −48           | +10      | <b>yes</b>           |
| Qwen3.5-27B                                       | −1        | −35           | +24      | <b>yes</b>           |
| GPT-5.4-mini                                      | −2        | −25           | +15      | <b>yes</b>           |
| GPT-5.4                                           | +0.3      | −36           | +25      | <b>yes</b>           |
| Claude-Sonnet-4-6                                 | −4        | −40           | +17      | <b>yes</b>           |
| Claude-Opus-4-6                                   | −5        | −17           | +9       | <b>yes</b>           |
| GPT-5.5                                           | −1        | −31           | −4       | no (all −)           |
| Claude-Opus-4-7                                   | −2        | −34           | −9       | no (all −)           |
| Sign-flip rate within protocol-controlled subset: |           |               |          | 12/15 = 80%          |




Figure 9: **Cross-task Gap sign matrix.** Each cell shows the per-model confidence–accuracy Gap (pp) for one of MMLU, GSM8K, HumanEval, BFCL; red = overconfident, blue = conservative. ★ on the right marks models with at least one sign reversal across the three same-protocol extension tasks (12/15 = 80%). The MMLU column is estimated from per-model NoMem–accuracy comparisons and is shown for context only; the sign-flip count uses only the three same-protocol tasks.




Figure 10: **MMLU-learned Platt scaling: cross-task transfer.** Number of models (out of 11 with usable Platt parameters from MMLU) on which the transferred recalibration helps, ties, or hurts on each target task. Helps on HumanEval (11/11) but hurts on GSM8K (7/11), confirming that “confidence shrinkage” learned from MMLU’s typical overconfidence harms already-conservative GSM8K models—per-task boundaries are structurally necessary, not merely convenient.

**Interpretation.** 12 of 15 models (80%) exhibit at least one Gap sign reversal across the three same-protocol extension tasks. The three exceptions are Mistral-7B (uniformly overconfident across all task types) and the two newest models GPT-5.5 and Claude-Opus-4-7 (uniformly conservative across the three extension tasks but still overconfident on MMLU per main text). The qualitative claim that “the same model exhibits qualitatively different calibration regimes across task types” survives the protocol-controlled restriction, ruling out the alternative explanation that the heterogeneity is purely a protocol artifact (under that hypothesis, sign-flip rate within the protocol-controlled subset should be near 0).

1081 **Caveat.** This subset analysis cannot rule out task-content confounds (math vs. code vs. tool-calling  
 1082 differ in skill demands beyond protocol) and we make no claim that the magnitudes of the per-model  
 1083 Gap differences are protocol-invariant; the sign-flip count is the only claim we defend with this  
 1084 slicing.

## 1085 Q.2 Lite protocol: T=100 vs T=500 ranking preservation

1086 To reduce adoption cost, we evaluate whether a 100-turn protocol preserves the ranking information  
 1087 of the full 500-turn evaluation. We compute NoMemory ECE and CBF at  $T=100$  vs.  $T=500$  across  
 1088 all models and measure rank correlation:

|      | Metric | Pearson $r$ (100 vs 500) | Spearman $\rho$ |
|------|--------|--------------------------|-----------------|
| 1089 | ECE    | 0.899                    | 0.886           |
|      | CBF    | 0.892                    | 0.689           |

1090 The  $T=100$  protocol preserves approximately 90% of the ranking information on ECE ( $r=0.90$ )  
 1091 and CBF ( $r=0.89$ ). We recommend the full 500-turn protocol for benchmark submissions and the  
 1092 100-turn “lite” protocol for rapid screening of new methods, with the caveat that CBF rank correlation  
 1093 ( $\rho=0.69$ ) is lower than ECE, suggesting that domain-level self-assessment requires longer interaction  
 1094 horizons to stabilize.

## 1095 R Extended Related Work Discussion

1096 This section provides the detailed citation-by-citation discussion that was compressed in the main  
 1097 text’s Related Work section. We retain it here for reviewers and readers seeking the full lineage.

1098 **LLM Self-Knowledge and Metacognition (extended).** Kadavath et al. [18] established that LLMs  
 1099 possess latent representations of their own correctness via P(True) probing. Ji-An et al. [17] showed  
 1100 that LLMs can monitor internal activations through in-context learning within a low-dimensional  
 1101 “metacognitive space.” Binder et al. [4] demonstrated behavior prediction via fine-tuning. Ackerman  
 1102 [1] evaluated LLM metacognition without self-reports, finding frontier models show limited, context-  
 1103 dependent metacognitive signals. Liu and van der Schaar [30] argued that truly self-improving  
 1104 agents require *intrinsic* metacognitive learning. Most directly relevant, Barkan et al. [3] found all 15  
 1105 tested LLMs systematically overconfident on BigCodeBench and SWE-Bench, with overconfidence  
 1106 worsening through multi-step tasks.

1107 **Knowledge Boundaries and Calibration Benchmarks (extended).** Li et al. [24] proposed a  
 1108 formal four-type taxonomy of LLM knowledge boundaries. Yin et al. [50] introduced the PGDC  
 1109 method via projected gradient descent with semantic constraints. SimpleQA [44] measures static  
 1110 factuality calibration on 4,326 short-answer questions. MMLU [15] and TruthfulQA [28] evaluate  
 1111 knowledge-level calibration in single-turn settings. HELM [27] provides multi-dimensional evaluation  
 1112 including ECE. Geng et al. [12] survey confidence estimation methods comprehensively.

1113 **LLM Agent Memory Systems (extended).** MemGPT [35] pioneered virtual memory hierarchies  
 1114 for agents. Mem0 [33], MemOS [26], RMM [41], and A-MEM [47] developed increasingly sophisti-  
 1115 cated memory architectures. MEM1 [54] proposed constant-memory consolidation for multi-turn  
 1116 agents. Sarukkai et al. [38] showed that agents storing successful trajectories as in-context examples  
 1117 improve ALFWorld success from 73% to 89%. MemoryAgentBench [16] provides the first system-  
 1118 atic benchmark for memory agents over four core competencies. None of these systems maintain  
 1119 endogenous per-domain accuracy state—the central novelty isolated by SelfBound.

1120 **Learning from Implicit Feedback (extended).** Tucker et al. [43] formalized coactive learning  
 1121 from user edits but assumes structured, reliable feedback. Liu et al. [29] analyzed WildChat and  
 1122 LMSYS dialogue logs, finding that over half of follow-up utterances contain feedback signals, but  
 1123 that retrospective distillation from such feedback yields mixed results—helping on short MTBench  
 1124 questions but not on longer WildBench ones. Leng et al. [23] showed that RLHF training itself  
 1125 induces systematic overconfidence.

**Selective Prediction and Abstention (extended).** Kapoor et al. [19] showed 1,000-sample LoRA fine-tuning produces generalizable uncertainty estimators. Diao et al. [9] (NAACL 2024 Outstanding Paper) demonstrated R-Tuning teaches abstention as a meta-skill. Xu et al. [46] trained models to generate self-reflective rationales. Traub et al. [42] (NeurIPS 2024 Spotlight) proposed AUGRC for selective classification. Li et al. [25] (TACL 2025) provides a comprehensive survey. Kirichenko et al. [20] introduced AbstentionBench, finding that reasoning fine-tuning degrades abstention by 24% on average. Ahdritz et al. [2] separated epistemic from aleatoric uncertainty using linear probes on frozen embeddings.

**Simulated Users and Closest Competitors (extended).** Zhou et al. [53] quantified the Sim2Real gap, finding LLM simulators create an “easy mode” with positivity bias. Chen et al. [5] introduced OmniBehavior from real user interaction logs, revealing models simulate users as an “active average person.” CCG [11] proves LLMs can adjust confidence through structured interaction but uses perfect binary feedback. LePe [14] converts self-consistency data into SFT labels but requires offline weight updates. SaySelf [46] optimizes per-query abstention without persistent memory. Dong et al. [10] formalized “capability boundary collapse” in RL-trained LLMs. Shen et al. [39], Chen et al. [6] (EMNLP 2024 Findings), and Mangala et al. [32] studied task-specific calibration variance.

## S API Model Versions

We list the exact model identifiers and access dates for all API-based models used in this benchmark:

| Model                    | Endpoint                 | Access Dates        |
|--------------------------|--------------------------|---------------------|
| Qwen-Turbo               | dashscope.aliyuncs.com   | 2026-04-10 to 04-12 |
| Qwen-Max                 | dashscope.aliyuncs.com   | 2026-04-10 to 04-12 |
| Qwen-Plus                | dashscope.aliyuncs.com   | 2026-04-10 to 04-12 |
| Qwen2.5-7B-instruct      | dashscope.aliyuncs.com   | 2026-04-10 to 04-12 |
| Qwen3.5-27B              | dashscope.aliyuncs.com   | 2026-04-10 to 04-12 |
| DeepSeek-V3.2            | dashscope.aliyuncs.com   | 2026-04-10 to 04-12 |
| GLM-5                    | dashscope.aliyuncs.com   | 2026-04-10 to 04-12 |
| Claude-Sonnet-4-6        | proxy / anthropic        | 2026-04-15          |
| Claude-Opus-4-6          | proxy / anthropic        | 2026-04-15          |
| GPT-5.4                  | proxy / openai           | 2026-04-15          |
| GPT-5.4-mini             | proxy / openai           | 2026-04-15          |
| Mistral-7B-Instruct-v0.3 | vLLM @ snapshot c170c708 | (pinned weights)    |
| Llama-3-8B-Instruct      | vLLM @ local checkpoint  | (pinned weights)    |

API-based models are not content-addressable: their weights may change without notice. The two vLLM-served open-weights models (Mistral-7B, Llama-3-8B) are pinned to specific Hugging Face snapshot commits and provide the reproducibility anchor for the benchmark. Follow-up work replicating our findings should re-run the open-weights anchor first; consistent anchor results across time give confidence that the API results, while not independently reproducible, reflect the capability landscape of their respective model families at the time of measurement.

## T MMLU Subject Selection

The 20 MMLU subjects used in SelfBound were selected to maximize diversity across STEM, humanities, and social sciences. The subjects and their per-model accuracy tiers are listed below:

## U User Persona Details

Each simulated user persona is defined by the following parameters:

- **Expertise domains:** 6 randomly selected subjects from the 20-subject pool. Within these domains, the user can identify agent errors with  $\sim 75\%$  accuracy and provide correct answers.
- **Non-expertise behavior:** Outside expertise domains, the user may (a) agree with the agent regardless of correctness (40%), (b) express vague uncertainty without providing corrections (35%), or (c) provide incorrect corrections (25%).

Table 7: MMLU subjects and per-model accuracy tiers (S = Strong, M = Medium, W = Weak).

| Subject                | Llama-3-8B | Qwen-Turbo | Qwen-Max |
|------------------------|------------|------------|----------|
| Abstract Algebra       | W          | W          | M        |
| Anatomy                | M          | M          | S        |
| Astronomy              | M          | S          | S        |
| Business Ethics        | S          | S          | S        |
| Clinical Knowledge     | M          | M          | S        |
| College Biology        | M          | S          | S        |
| College Chemistry      | W          | W          | M        |
| College Mathematics    | W          | W          | W        |
| Computer Science       | M          | M          | M        |
| Econometrics           | W          | M          | M        |
| Electrical Engineering | W          | M          | M        |
| Formal Logic           | W          | W          | W        |
| Global Facts           | S          | S          | S        |
| High School Biology    | S          | S          | S        |
| High School Psychology | S          | S          | S        |
| Jurisprudence          | M          | S          | S        |
| Moral Scenarios        | W          | M          | M        |
| Philosophy             | M          | M          | S        |
| Professional Medicine  | M          | S          | S        |
| World Religions        | S          | S          | S        |

1161 • **Communication style:** Each persona is assigned one of three styles—*direct* (“No, the answer is  
1162 B”), *hedging* (“I think that might not be right, maybe it’s B?”), or *Socratic* (“Are you sure? What  
1163 about option B?”).

1164 The user simulator prompt includes the persona specification, the ground-truth answer, the agent’s an-  
1165 swer, and instructions to generate feedback consistent with the persona’s expertise and communication  
1166 style.

## 1167 V Baseline Implementation Details

1168 **B1: NoMemory.** The agent is prompted to provide a confidence score between 0 and 1 alongside  
1169 its answer. No history is provided beyond the current question.

1170 **B2: OraclePDC.** After each turn, the agent receives a binary correctness signal. It maintains a  
1171 dictionary mapping inferred domain labels to running accuracy statistics (correct count / total count  
1172 per domain). The domain label is inferred by the agent from question content.

1173 **B3: User-Feedback PDC.** Same as B2, but the correctness signal is extracted from user feedback  
1174 by an LLM-based classifier (Qwen-Max prompted with the (Q, A, feedback) triple to output a  
1175 binary correctness judgment). LLM-based decoding gives near-uniform coverage across in- and  
1176 out-of-expertise feedback, removing rule-classifier coverage as a confound.

1177 **B6: LSA.** The agent receives a formatted summary of its 50 in-session interactions (question  
1178 subject, confidence, decoded correctness) and is prompted at the *end of the session* to output an  
1179 estimate of its accuracy on each domain.

1180 **B10: EM-Joint.** Joint estimation of  $(\hat{p}(d), \hat{r}_{\text{in}}, \hat{r}_{\text{out}})$  via Expectation–Maximization on  
1181 the (decoded feedback, raw confidence) stream. E-step:  $P(\text{correct}_t | f_t, c_t, p, r) \propto$   
1182  $P(f_t | \text{correct}_t, r) P(\text{correct}_t | c_t, p)$ . M-step:  $\hat{p}(d) \leftarrow$  posterior mean of correctness in domain  $d$ ;  
1183  $\hat{r}_{\text{in/out}} \leftarrow$  posterior agreement between feedback and inferred correctness within each expertise group.

1184 **B11: ZS-SK.** A single prompt at the start of evaluation: “For each of the 20 MMLU subjects,  
1185 estimate your own accuracy on a 0–100 scale, with no access to any prior interaction.” The estimates  
1186 are parsed into  $\hat{p}(d)$  and CBF is computed once.

1187 **B12: GT-SelfStats.** The true per-domain accuracy  $p^*(d)$  for the agent (from the open-loop pre-  
1188 test) is provided in the prompt and the agent is asked to repeat each value. Diagnostic anchor for  
1189 verbalization.

## 1190 W Extended Results

Table 8: Per-tier CBF breakdown for representative models. Across baselines, CBF is highest for strong domains and lowest for weak domains, indicating that agents are systematically better at recognizing their strengths than their weaknesses.

| Model      | Baseline | Strong | Medium | Weak |
|------------|----------|--------|--------|------|
| Llama-3-8B | ZS-SK    | 0.79   | 0.68   | 0.51 |
|            | LSA      | 0.76   | 0.63   | 0.50 |
| Qwen-Turbo | NoMemory | 0.81   | 0.70   | 0.55 |
|            | ZS-SK    | 0.83   | 0.72   | 0.57 |
| Qwen-Max   | NoMemory | 0.80   | 0.71   | 0.55 |
|            | LSA      | 0.85   | 0.75   | 0.58 |

## 1191 X Frontier Sub-Sample: Sample-Size Sensitivity

1192 The frontier-only analyses use  $n = 6$  models. We compared the frontier within-tier means under  
1193 two session budgets to verify robustness: a 10-session-per-model lite protocol and the 50-session-  
1194 per-model main protocol used throughout the paper. The qualitative pattern (LSA leads on frontier;  
1195 ZS exceeds NoMem on GPT but falls below NoMem on Claude) is preserved under both. The  
1196 per-model magnitudes shift most on the Claude rows (ZS-SK falls 0.71–0.81  $\rightarrow$  0.46–0.48 with the  
1197 larger sample; LSA shifts 0.66–0.77  $\rightarrow$  0.64–0.72), confirming that the smaller sample inflated ZS  
1198 estimates on Claude due to chance alignment with limited per-domain accuracy estimates. On the  
1199 three GPT rows the changes are smaller ( $\leq 0.05$  in any cell). The 50-session estimates are reported  
1200 throughout the main paper.

## 1201 Y Reproducibility Checklist

1202 We follow the NeurIPS 2026 Datasets & Benchmarks reproducibility template. Cross-references  
1203 point to specific paper sections.

### 1204 Claims and contributions.

- 1205 • *All claims explicitly stated.* ✓ Abstract: 4 explicit claims (i–iv); Introduction: 3 hypotheses  
1206 framed as testable (H1/H2/H3).
- 1207 • *Limitations discussed.* ✓ §5 (10 paragraphs): question format, cross-task protocol confound,  
1208 held-out controls, evaluation cost, simulator validity, model coverage, session length, and explicit  
1209 research-direction sketches.
- 1210 • *Negative results reported.* ✓ R4 LOOCV minimal overfit gap of +0.003 — confirms the  
1211 parameter-free  $\theta = 0$  rule generalizes (Appendix E); R4 misses the GPT-5.5 case where LSA >  
1212 ZS (§4.3); LSA catastrophic failure on Qwen-Plus (−0.300 from ZS, Appendix E); first-order  
1213 single-rater methods (UF-PDC, Platt, HistBin, Bayes, EM-Joint) collapse near NoMem on frontier  
1214 (§E).

### 1215 Data and code.

- 1216 • *Code released.* ✓ the anonymized repository: scheduler, baseline implementations (10), DS /  
1217 GLAD scripts, R4 diagnostic / analysis tool, ablation scripts.
- 1218 • *Data released.* ✓ All 30k+ session traces (raw  $(q_t, a_t, c_t, f_t, \text{user}_t, \text{decoded}_t)$  tuples) at the same  
1219 repo, organized per model  $\times$  baseline.
- 1220 • *Pinned model versions.* ✓ §S: open-weights models pinned to HuggingFace snapshot commits;  
1221 API models reported with exact endpoints, identifiers, and access dates.

- *Random seeds.* Sessions are deterministic given (model, persona, question pool); we fix question shuffling per (model, session) with a 42-derived seed; user-persona temperature is fixed at 0.7 (documented in Appendix U).

## Experimental protocol.

- *Train/validation/test splits.* ✓ Per-domain held-out split for  $p^*(d)$  ground truth (§3); MMLU-Pro and BBH held-out evaluation for contamination control (§4.2).
- *Hyperparameters.* ✓ All baseline hyperparameters listed in Appendix C: episodic memory  $k = 5$ , EMA  $\alpha = 0.3$ , EM iterations 50, GLAD/DS init values, etc.
- *Statistical significance.* ✓ Bootstrap  $200\times$  session-resampled SDs (Appendix G); pairwise sign tests with explicit  $p$ -values; Bonferroni correction discussed where applicable.
- *Compute resources.* §5: total run cost  $\sim$  \$300–\$500 in commercial-API calls plus open-weights compute; protocol evaluable on a single A100 / RTX A6000 workstation.

## Reproducibility risks and mitigations.

- *API model deprecation.* §S (extended manifest): per-model risk classification (HuggingFace-pinned, API-snapshot, API-rolling); two open-weights models (Mistral-7B, Llama-3-8B) constitute the time-stable replication anchor.
- *Simulator stability.* The user simulator is itself an LLM; we use cross-family pairings to control for dual-role bias (§M, §N) and report a cross-simulator ablation showing  $\Delta\text{CBF} < 0.005$  across alternative simulator choices.
- *Bug-fix disclosure.* ✓ §A: prior versions (pre-fix) silently fell back algorithmic L2 baselines to LLM verbal estimates; the patch and full re-run are documented; pre-patch values are not used in any quantitative claim.
- *Adaptive policy R4 LOOCV gap.* ✓ R4 uses a fixed threshold  $\theta = 0$  by design (parameter-free); explicit LOOCV shows the gap from in-sample (0.722) to held-out (0.719) is only  $+0.003$ , confirming generalization (Appendix E).

## Ethics, broader impact, and licensing.

- *No human subjects in the released benchmark.* All “user” interactions are LLM-simulated personas; no human data collection. A planned future human-validation pilot study (50 sessions, 3 annotators) will require IRB approval.
- *Broader impact.* The benchmark exposes failure modes (LSA on GPT-5.5/Qwen-Plus; “feedback as friend or foe” split) directly relevant to deploying LLM agents that claim self-knowledge in production; we hope it informs deployment caution rather than enabling new harms.
- *Licensing.* SelfBound released under MIT (code) and CC-BY 4.0 (data/traces), consistent with the Datasheet license. MMLU usage follows the original MMLU license.